

Advanced Computer Networks Lab



Dr Sudipta Saha

Associate Professor

Dept of Computer Science & Engineering

Indian Institute of Technology Bhubaneswar

Tutorial-2

Teaching Assistant: Nitin



**DSSRG: Decentralized
Smart Systems Research
Group**

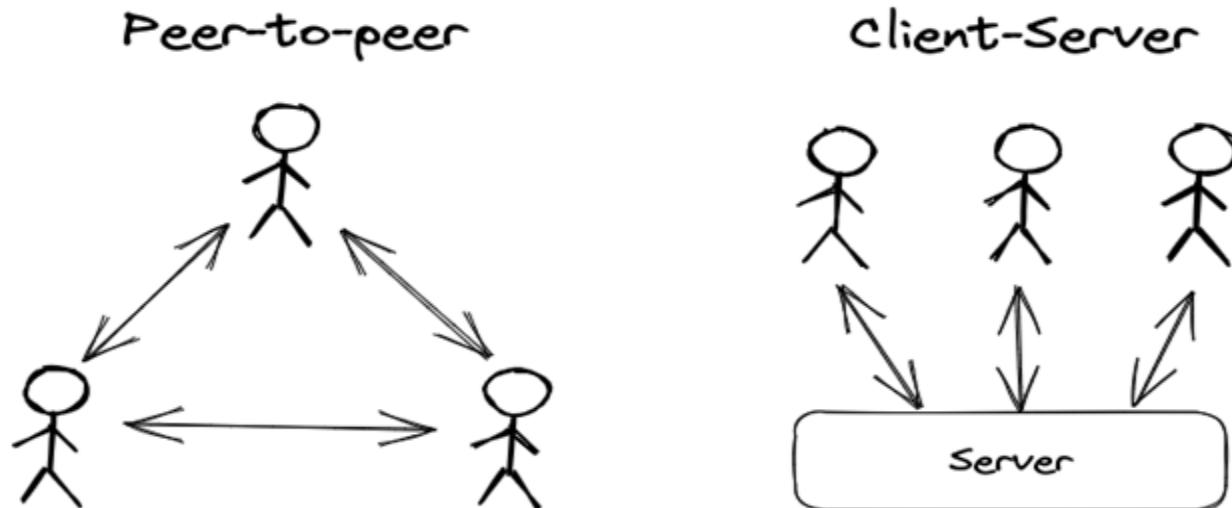
<https://sites.google.com/iitbbs.ac.in/dssrg>
Or Google dssrg iitbbs



p2p Networks

A decentralized network model where all devices, or "peers," have equal status. Each computer can act as both a client (requesting resources) and a server (sharing resources) directly with others, without needing a central authority.

- Client-Server: Centralized. One powerful server provides services to many clients.
- P2P: Distributed. All peers are equal, sharing resources directly among themselves.

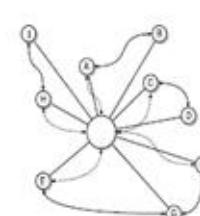


p2p Unstructured Topologies

A decentralized network design, common in P2P systems, where connections between peers are established arbitrarily and without a specific organizational pattern. There is no predefined rule or map for how the network is formed.

Key Characteristics:

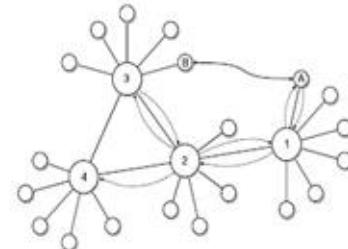
- Random Connections: Peers connect to whichever neighbors they discover, often forming random, ad-hoc links.
- Search by Flooding: To find a file or resource, queries are broadcast (flooded) to all connected neighbors, which is simple but can generate significant network traffic.
- Resilient & Easy to Join: Highly robust to peers frequently joining/leaving (churn) and easy for new peers to integrate.
- Inefficient Search: There is no guarantee of locating a resource, even if it exists in the network.



(a) Napster & BitTorrent



(b) Gnutella



(c) Gnutella/Overnet/eDonkey2000

Common Use: Early file-sharing networks
(e.g., Gnutella), Napster, bittorrent.



p2p Structured Topologies

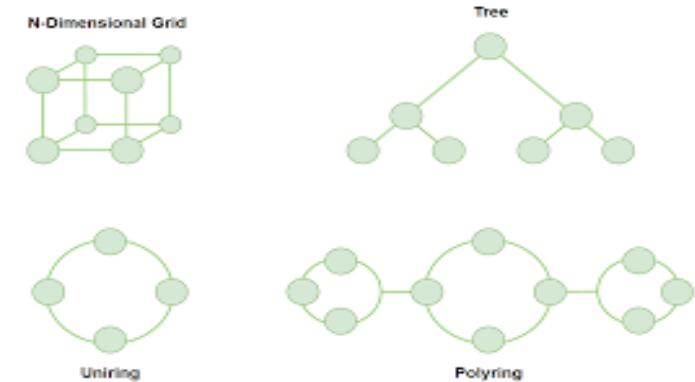
A controlled P2P network architecture where connections between peers are organized according to a specific, deterministic rule set—most commonly a Distributed Hash Table (DHT). This creates a predictable global structure that allows efficient resource location.



p2p Structured Topologies

Key Characteristics:

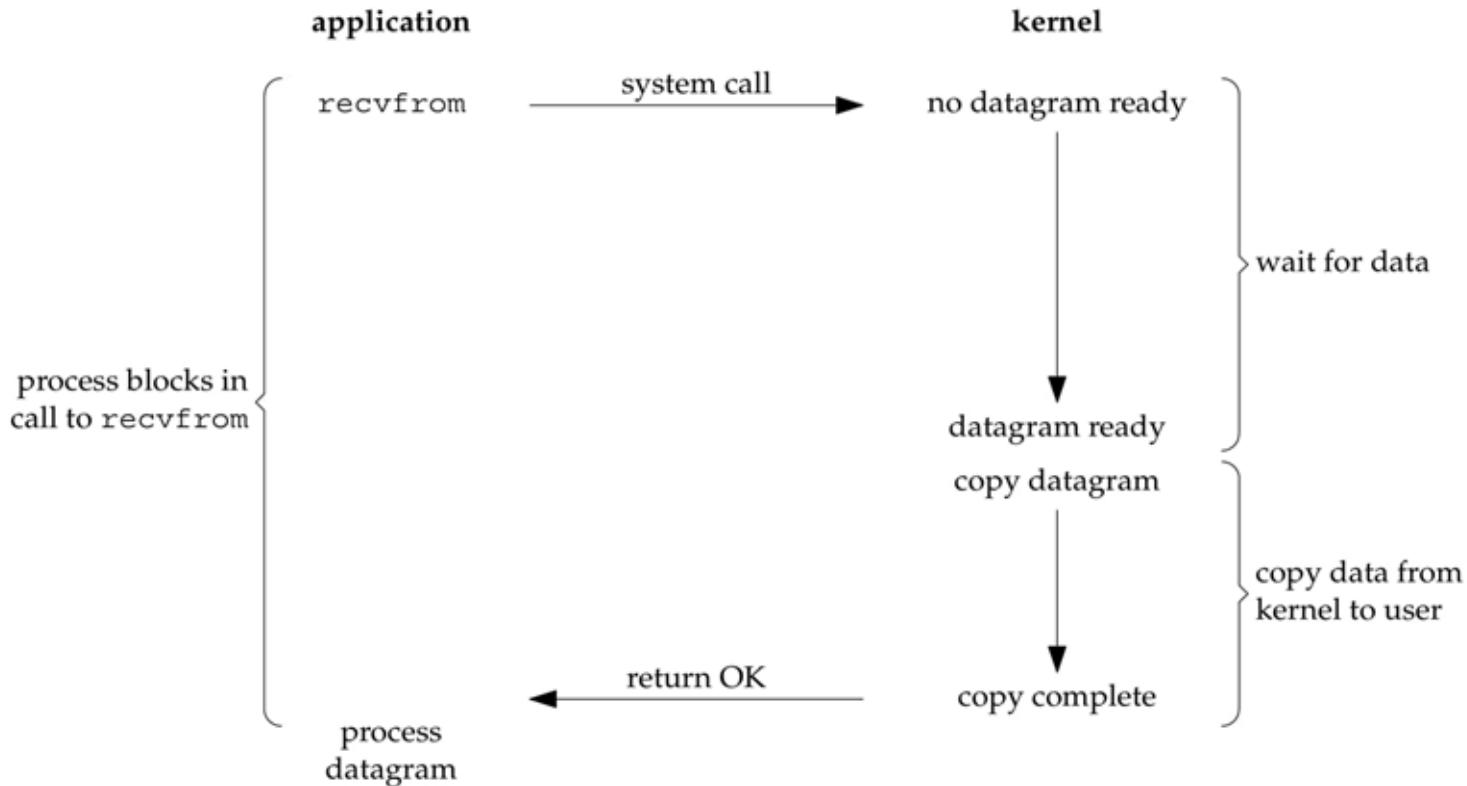
- Organized & Predictable: Peers and the resources they hold are mapped to specific, logical positions within the network using unique identifiers (like keys).
- Efficient Lookup: Any resource can be reliably located within a small, predictable number of steps (usually $O(\log n)$), without flooding the network.
- Scalable: Provides guaranteed discovery and scales efficiently to very large networks.
- Maintenance Overhead: Requires more complex protocols to maintain the structure as peers join and leave (churn).



BLOCKING MODEL

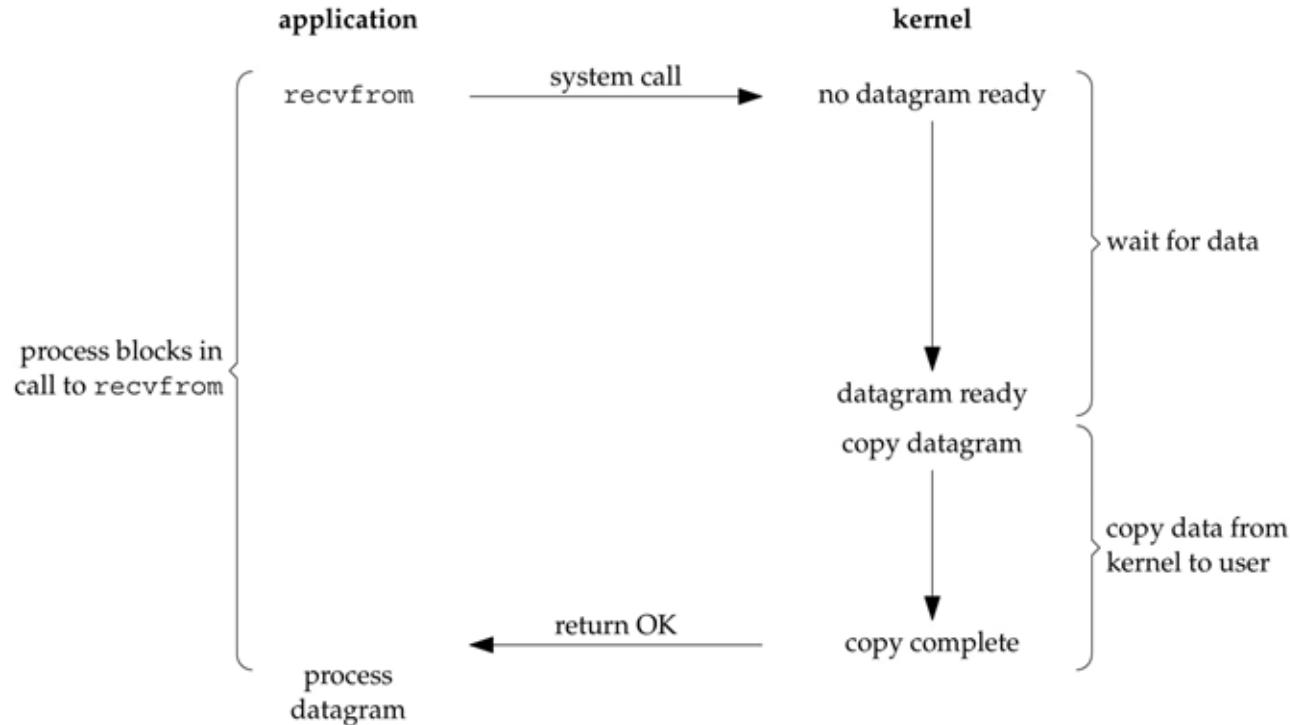
A default mode of operation in socket programming where a function call (like `send()` or `recv()`) halts the execution of the program until the requested network operation is fully complete or an error occurs.

Blocking (Synchronous): "Wait here until this specific I/O operation is done."



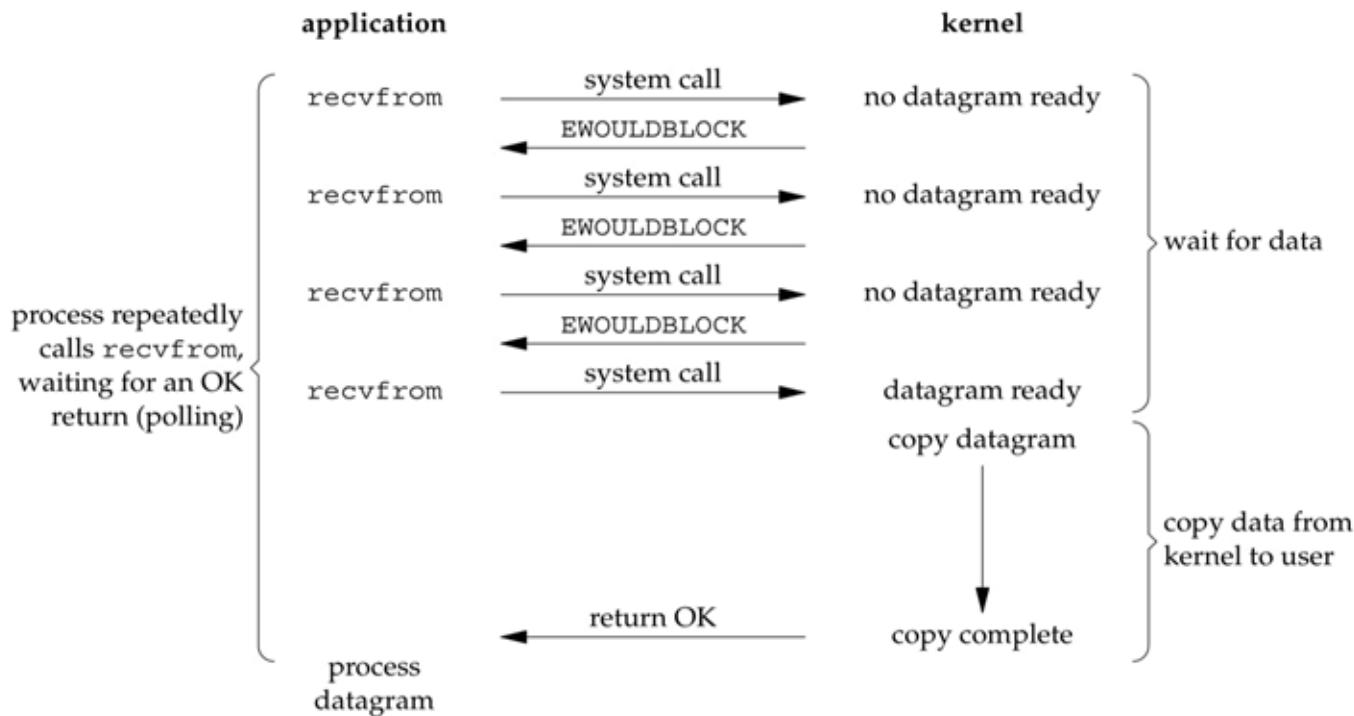
BLOCKING MODEL

- `accept()` - Sleeps until someone connects
- `connect()` - Sleeps until connection succeeds/fails
- `recv()` - Sleeps until data arrives
- `send()` - Sleeps (rarely) if buffer is full



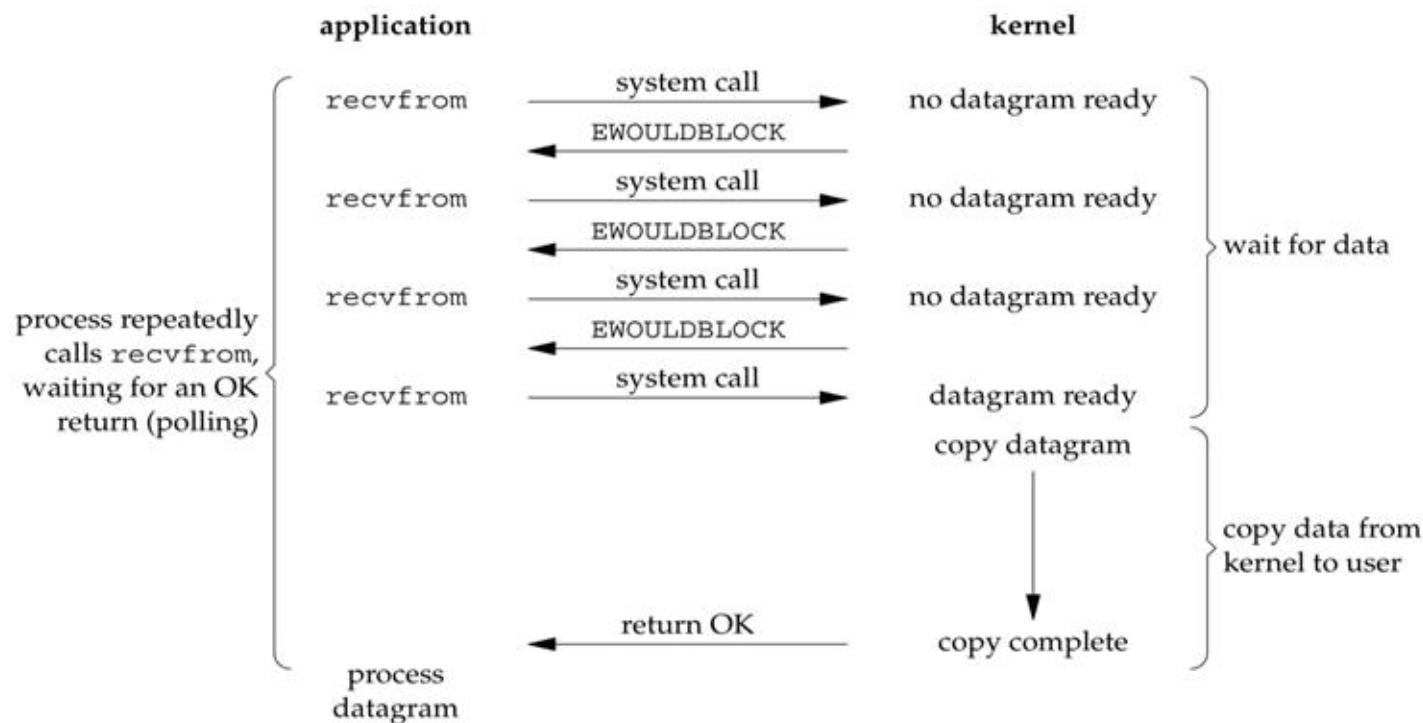
NON BLOCKING MODEL

An operational mode in socket programming where a function call (like `send()` or `recv()`) returns immediately, regardless of whether the operation could be completed. This allows the program to continue execution without waiting for the underlying network I/O to finish.



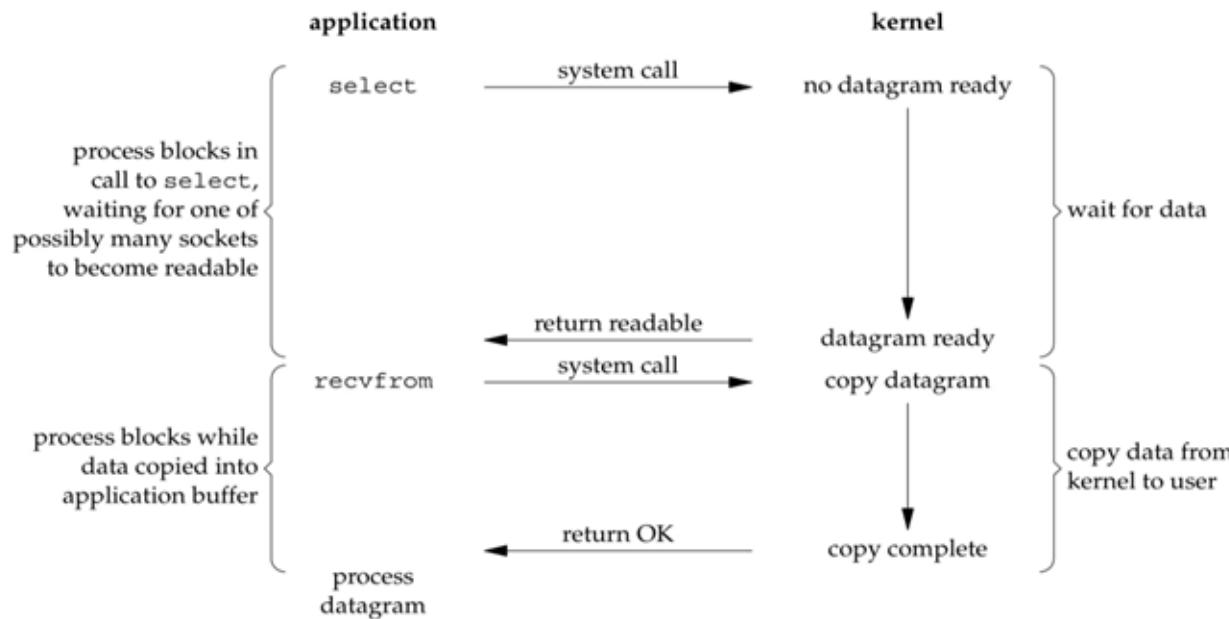
NON BLOCKING MODEL: POLLING

The technique of repeatedly checking the status of one or more non-blocking sockets to see if they are ready for a read or write operation. Instead of relying on the operating system to notify the program when data is available, the program actively and periodically inquires about each socket's state.



I/O MULTIPLEXING

A high-efficiency programming pattern that allows a single thread to monitor multiple socket descriptors simultaneously, and be notified when one or more become ready for I/O operations (e.g., readable or writable). It is the optimized implementation of polling for production systems.



select()

- A core function for I/O multiplexing that allows a program to monitor multiple file descriptors (typically sockets) simultaneously to see if any become ready for I/O operations. It's the original and most portable method for implementing event-driven network servers.
- The program passes three sets of file descriptors to select(): those it wants to monitor for read readiness, write readiness, and exceptional conditions. The call blocks until at least one descriptor in any set becomes ready or until a timeout expires.
- Descriptor Sets: Uses fixed-size bitmask arrays (fd_set) to represent sets of descriptors, limiting the maximum number of descriptors that can be monitored (typically 1024).
- Linear Scanning: On return, the program must iterate through all monitored descriptors to determine which ones are ready, which becomes inefficient with many connections.
- Destructive Call: The sets are modified by the call; they must be rebuilt before each subsequent select() invocation.



select()

```
int select(int maxfdp1, fd_set *readset, fd_set *writeset, fd_set *exceptset, const struct timeval *timeout);
```

Returns: positive count of ready descriptors, 0 on timeout, -1 on error.

- **maxfdp1**:The maximum file descriptor value + 1 in all sets. Determines how much of the descriptor space the kernel scans.
- **readset**:Set of descriptors to monitor for read readiness (incoming data, connection requests, socket closure).
- **writeset**:Set of descriptors to monitor for write readiness (when data can be sent without blocking).
- **exceptset**:Set of descriptors to monitor for exceptional conditions (out-of-band data, socket errors).
- **timeout**:Maximum time to wait for activity. NULL = block forever, 0 = return immediately (poll), non-zero = wait specified time.



select()

- Each bit position represents a file descriptor number
- Example: Bit 5 = descriptor 5, Bit 32 = descriptor 32
- One fd_set can track hundreds of descriptors in just a few bytes.
- Defined by FD_SETSIZE (typically 1024 on most systems), Cannot monitor descriptors ≥ 1024 without recompiling kernel, Wastes memory if monitoring only a few descriptors
- After select() returns, must check all possible descriptors 0...maxfd
- Inefficient for checking a few active descriptors among many possible

```
void FD_ZERO(fd_set *fdset); /* clear all bits in fdset */
```

```
void FD_SET(int fd, fd_set *fdset); /* turn on the bit for fd in fdset */
```

```
void FD_CLR(int fd, fd_set *fdset); /* turn off the bit for fd in fdset */
```

```
int FD_ISSET(int fd, fd_set *fdset); /* is the bit for fd on in fdset ? */
```



poll()

An alternative I/O multiplexing function that overcomes the fixed descriptor limit of select(). Instead of bitmasks, poll() uses an array of structures to monitor file descriptors.

```
int poll(struct pollfd *fds, nfds_t nfds, int timeout);
```

Returns: count of ready descriptors, 0 on timeout, -1 on error

```
struct pollfd {  
    int fd;      /* file descriptor */  
    short events; /* events to watch for (input) */  
    short revents; /* events that occurred (output) */  
};
```

- **nfds_t nfds** specifies the number of elements in the **fds** array that should be processed by **poll()**. It tells the kernel: "Check only the first nfds entries in this array."
- **timeout:** Maximum time to wait for activity. **NULL** = block forever, **0** = return immediately (**poll**), non-zero = wait specified time.



poll(): Common Event Flags

Flag	Meaning	When Used
POLLIN	Data available to read	Normal/priority data ready
POLLPRI	Urgent data available	Out-of-band data on TCP socket
POLLOUT	Writing will not block	Ready for output
POLLERR	Error condition	Always monitored automatically
POLLHUP	Hang up	Connection closed
POLLNVAL	Invalid request	Descriptor not open



select() vs poll()

Feature	select()	poll()
Descriptor Limit	Fixed (FD_SETSIZE, typically 1024)	Dynamic (limited by system resources)
Event Separation	Input/Output sets modified on return	Separate events (input) and revents (output) fields
Portability	Widely portable	Almost as portable (some older Unix systems may not have it)
Performance	$O(n)$ scan of all watched descriptors	$O(n)$ scan of array elements



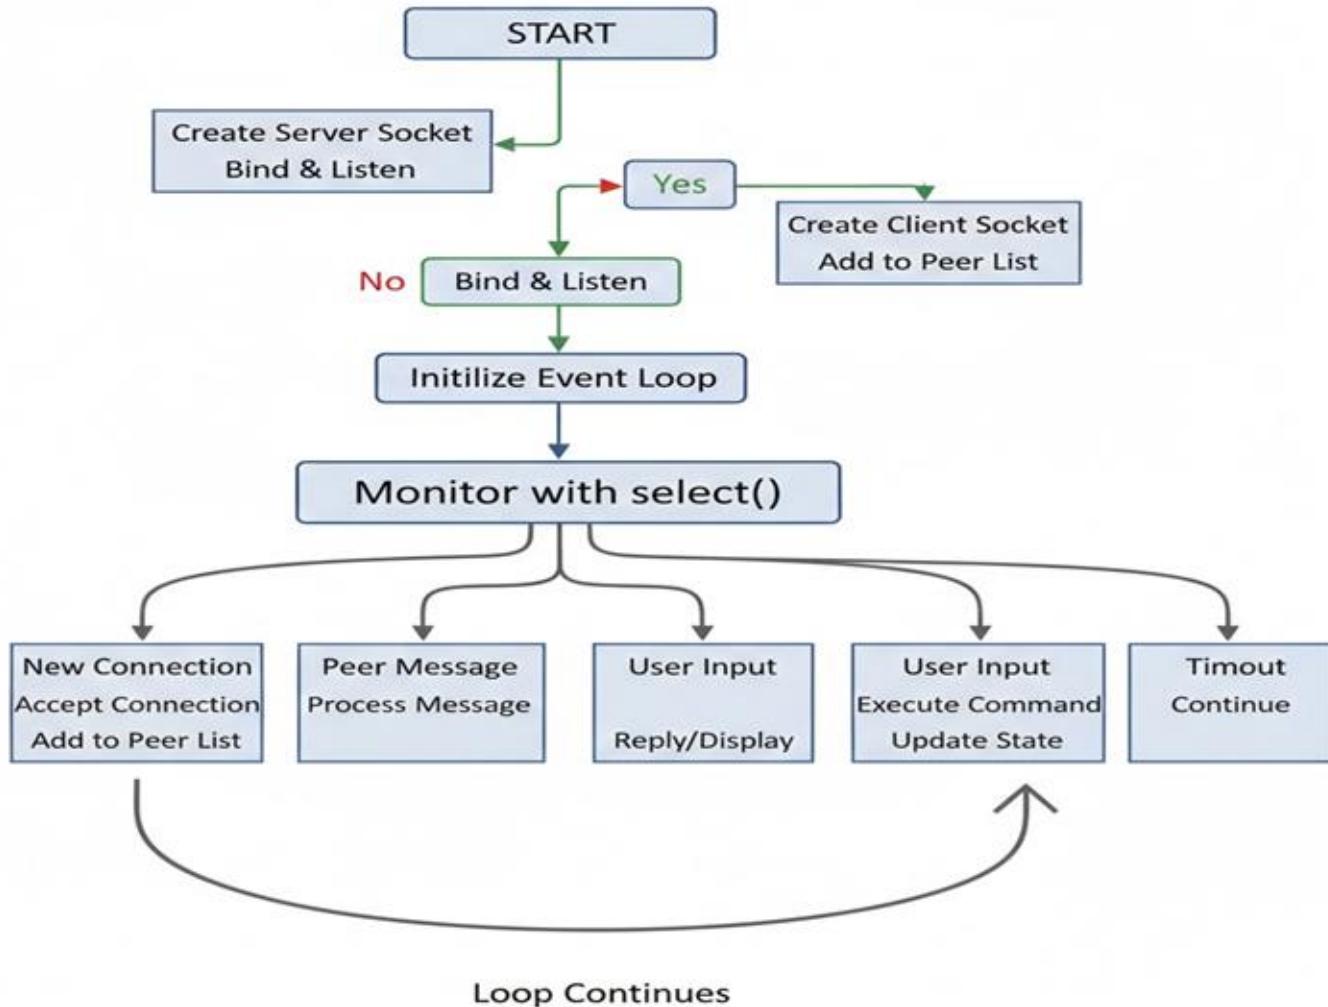
Writing a single threaded peer code

```
int main() {  
    // 1. Initialize peer socket  
    peer_socket = setup_peer();  
    // 2. Add initial connections  
    add_to_monitor(peer_socket, POLLIN);  
    add_to_monitor(STDIN_FILENO, POLLIN); // User input  
    while (running) {  
        // 3. Wait for ANY activity (connections, messages, input)  
        int ready = poll(monitor_set, num_monitored, -1);  
        // 4. Process all ready events  
        for (each monitored descriptor) {  
            if (descriptor == STDIN) {handle_user_command();}  
            else if (descriptor == peer_socket) {  
                handle_incoming_connection();}  
            else if (descriptor is ready for reading) { process_message(descriptor);}  
            else if (descriptor is ready for writing) {end_pending_data(descriptor);}  
        }  
    }  
}
```



Writing a single threaded peer code

P2P Node Event Loop Architecture



QUERYING

In Peer-to-Peer (P2P) networks, there is no central server to manage data. Therefore, finding a specific file or resource requires decentralized search mechanisms. The efficiency of a query depends heavily on the network's **topology** (Structured vs. Unstructured).

Flooding	Unstructured	A node sends a query to all neighbors, who forward it until the TTL (Time-to-Live) expires.	Pros: Simple. Cons: High bandwidth overhead (broadcast storm).
Random Walk	Unstructured	A node sends a query to a limited number of random neighbors.	Pros: Lower traffic than flooding. Cons: Slower discovery time.
DHT Lookup	Structured	Uses a Distributed Hash Table (e.g., Chord, Kademlia) to map keys to specific nodes.	Pros: Extremely efficient ($O(\log n)$ hops). Cons: High maintenance



Flooding

Flooding is the fundamental search algorithm used in unstructured P2P networks (e.g., Gnutella v0.4). It operates on a "broadcast" principle.

- 1. Query Generation:** The source node (S) creates a query descriptor containing:
 - **Search Criteria:** Keywords or metadata.
 - **Message ID (GUID):** A unique identifier to track the query.
 - **TTL (Time-to-Live):** A counter (typically set to 5–7) to limit propagation range.
- 2. Propagation (The Ripple Effect):**
 - Node S sends the query to *all* of its directly connected neighbors.
 - Each neighbor receives the query, decrements the TTL by 1, and forwards it to *all* of its own neighbors (excluding the one it received it from).
 - This continues until TTL = 0.
- 3. Loop Detection:** Nodes track Message IDs they have seen recently. If a node receives a duplicate query (same ID), it discards it immediately to prevent infinite cycles.



References

- Beej's Guide to Network Programming Using Internet Sockets
Brian "Beej Jorgensen" Hall v3.3.0, Copyright © January 19, 2026
- Xuemin Shen · Heather Yu · John Buford · Mursalin Akon
Editors Handbook of Peer-to-Peer Networking Springer

