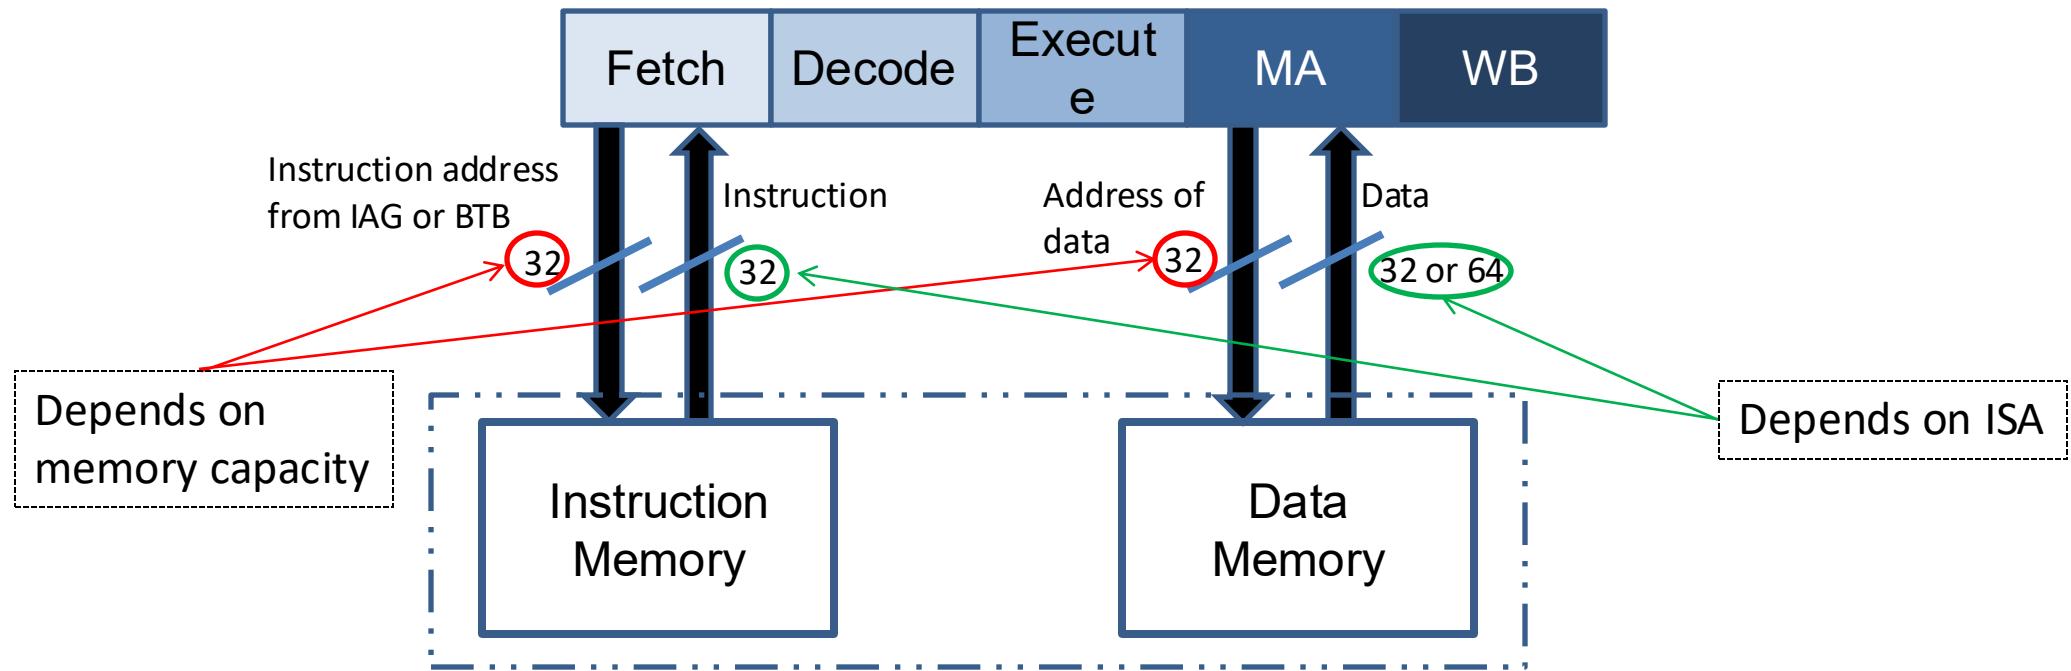


# **CS6L047: Advanced Computer Architecture**

## Cache

(Appendix B, Chapter 2, HP 6e)



- We usually treat “Instruction Memory” and “Data Memory” as black boxes so far.
- In this module, we’ll dive in!

# Ideal Memory!

- › **Fast** Speed matches with the processor
- › **Large** If it were up to us, we would wish for infinite memory space!
- › **Low cost** Obviously! If not free ☺
- › **Highly reliable** My data is precious to me! Stolen or lost!
- › **High Bandwidth** Infinite! Accessing one data or multiple should take same time.
- › **Consumes less power** Increasing interest in this aspect!
- › **Requires less area** Not so pressing at present!
- › **How to meet these requirements?**

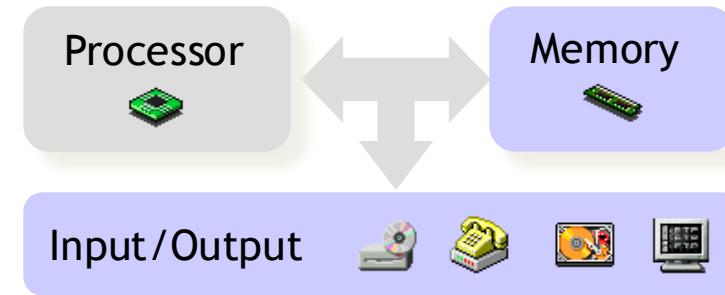


5MB hard disk in 1956 by IBM.



# Memory Systems

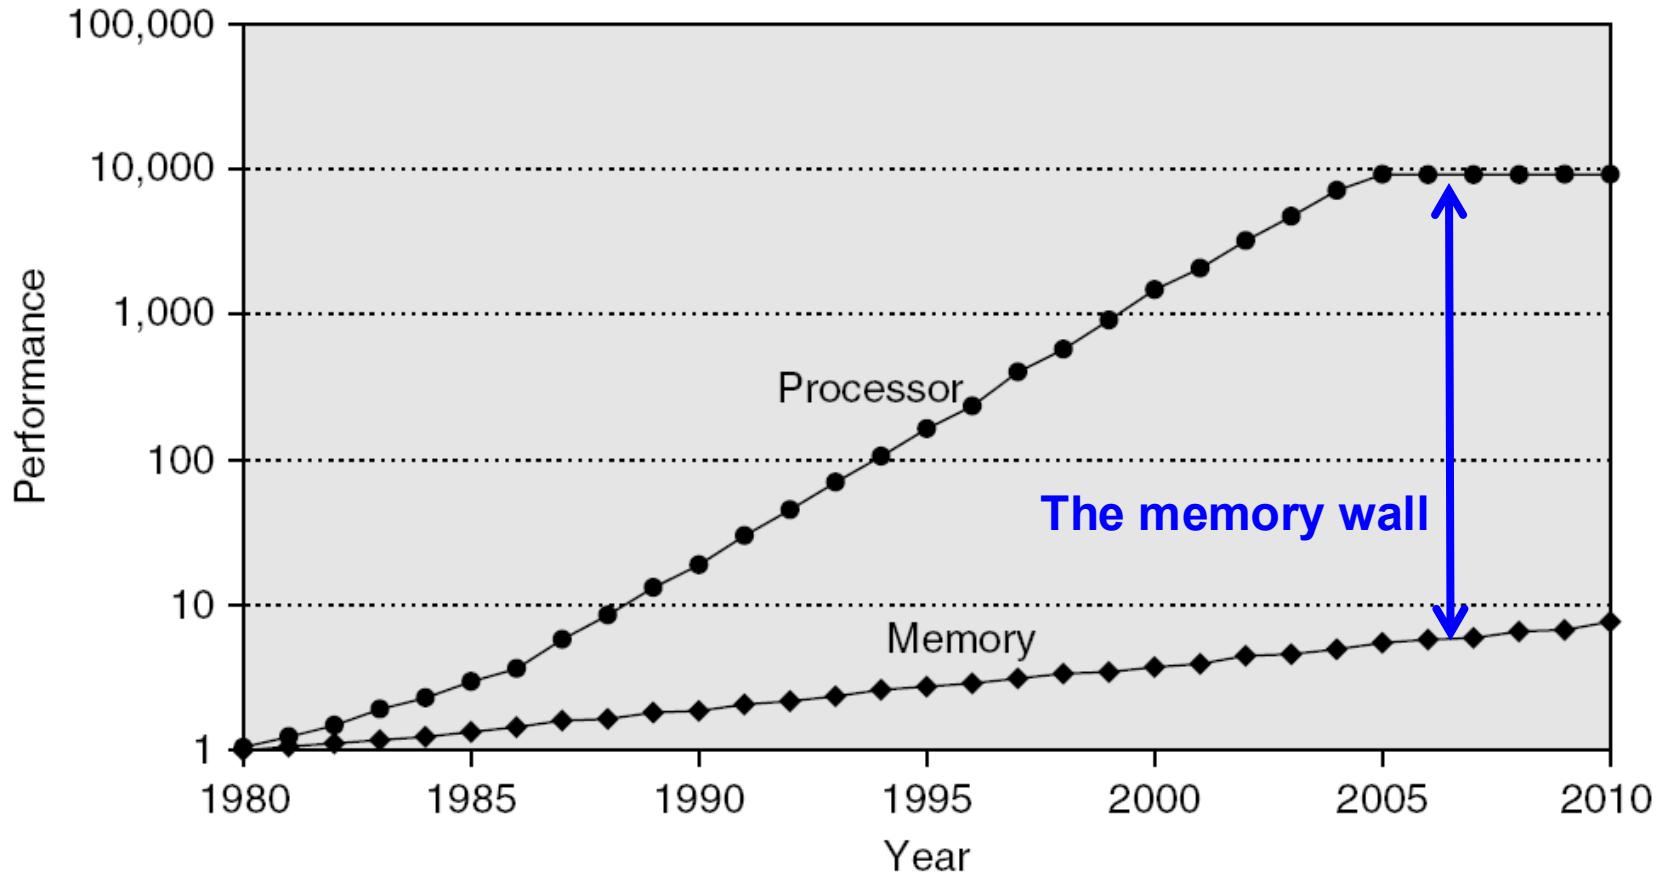
- › How can we supply the CPU with enough data to keep it busy?
- › We will focus on **memory** issues,
  - › which are frequently bottlenecks that limit the performance of a system.



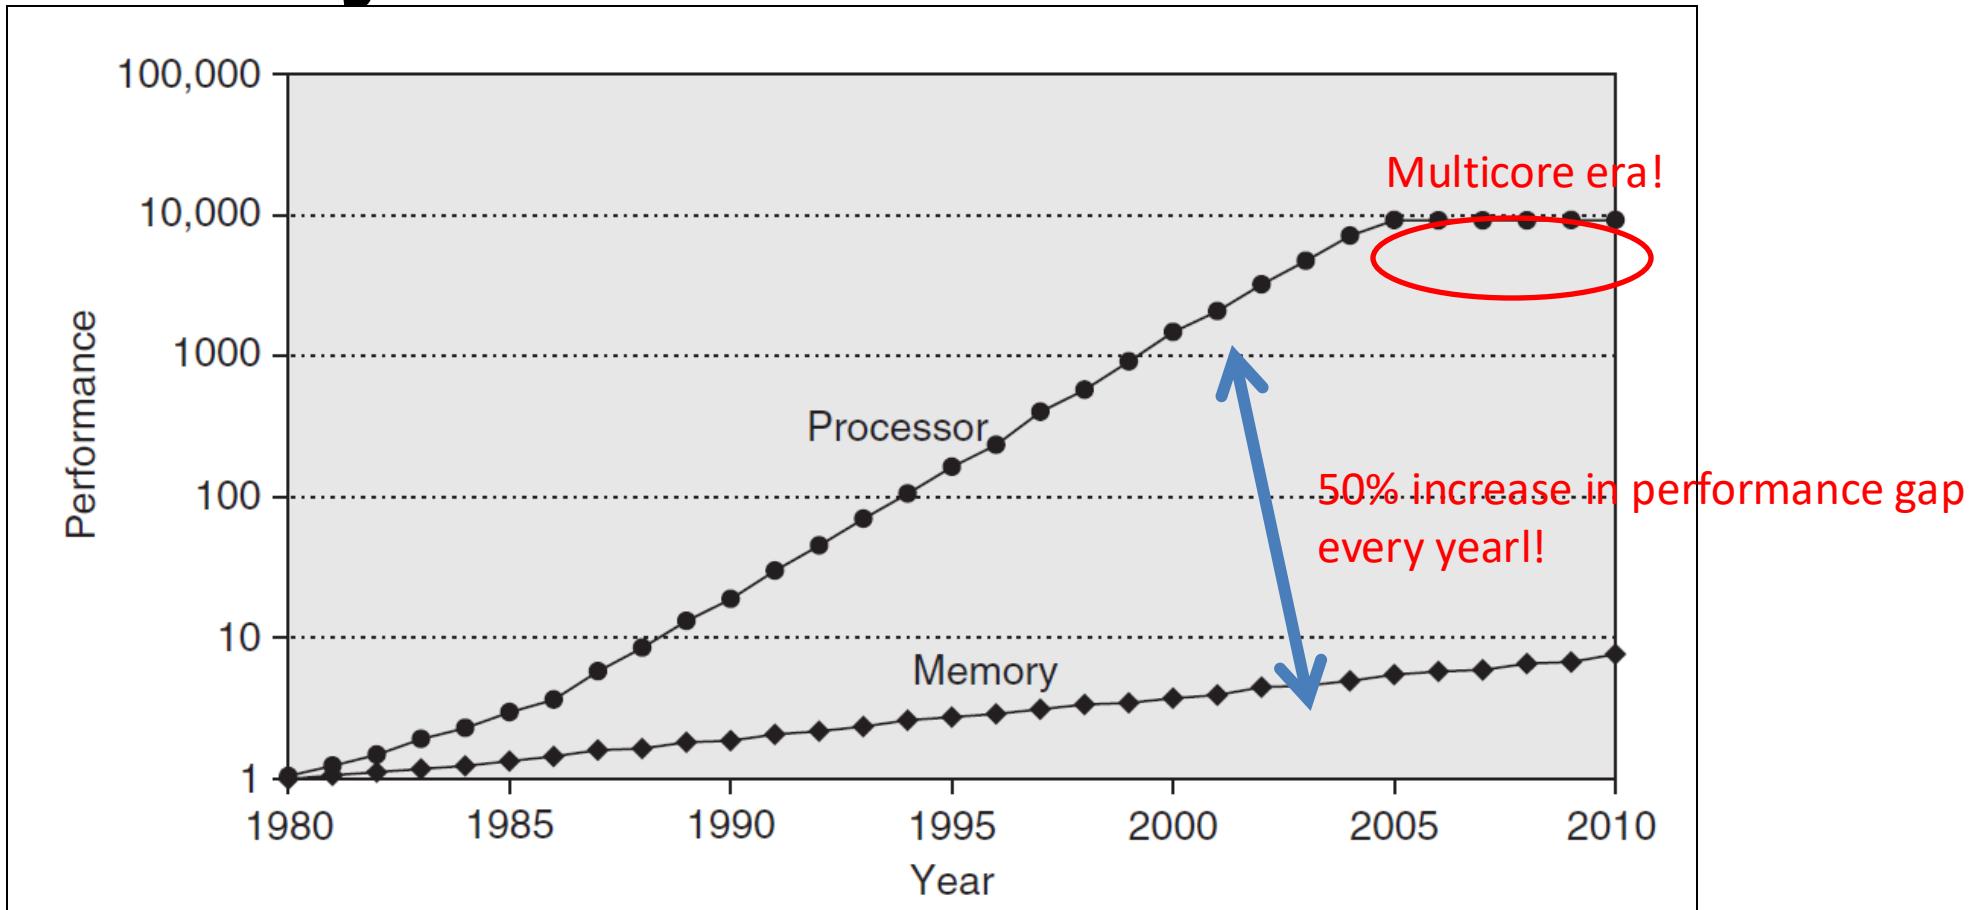
Storage	Speed	Cost	Capacity	Delay	Cost/GB
Static RAM	Fastest	Expensive	Smallest	0.5 – 2.5 ns	\$1,000's
Dynamic RAM	Slow	Cheap	Large	50 – 70 ns	\$10's
Hard disks	Slowest	Cheapest	Largest	5 – 20 ms	\$0.1's

- › Ideal memory: **large, fast and cheap**

# Performance Gap



# The Memory Wall



- Memory latency was an issue !
- Now, both latency and bandwidth are important !

From, Computer Architecture: A Quantitative Approach 4/e by Hennessey and Patterson

10-Nov-25

The Memory Hierarchy

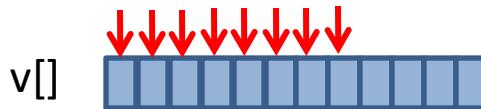
# Reality - A classic example

- You are sitting in the Library (preparing for end-sem):
  - Page you are reading
  - Book you are holding
  - Table you are sitting at, with possibly few books of current interest
  - Book racks
- What is the key observation here?
  - At any given moment, we would need only few books to read!
- Something similar with the programs we write –
  - Programs access a relatively small portion of their address space at any instant of time!  
Principle of locality!
  - Predominantly there are two types of locality:
    1. If a data location is referenced, it will tend to be referenced again in near future! Temporal locality!
    2. If a data location is referenced, nearby locations tend to get referenced soon! Spatial locality!

# Understanding Locality!

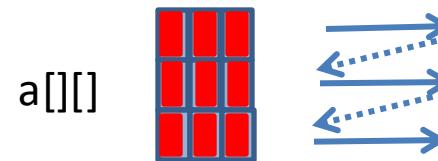
Ex1:

```
1 int sumvec(int v[N])
2 {
3     int i, sum = 0;
4
5     for (i = 0; i < N; i++)
6         sum += v[i];
7     return sum;
8 }
```



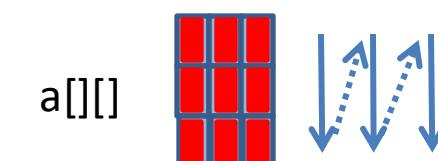
Ex2:

```
1 int sumarrayrows(int a[M][N])
2 {
3     int i, j, sum = 0;
4
5     for (i = 0; i < M; i++)
6         for (j = 0; j < N; j++)
7             sum += a[i][j];
8     return sum;
9 }
```



Ex3:

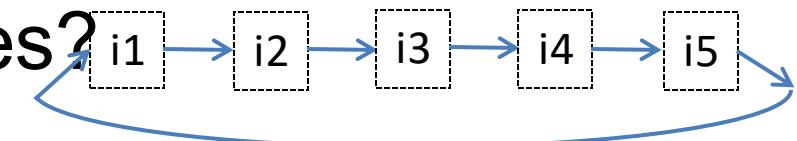
```
1 int sumarraycols(int a[M][N])
2 {
3     int i, j, sum = 0;
4
5     for (j = 0; j < N; j++)
6         for (i = 0; i < M; i++)
7             sum += a[i][j];
8     return sum;
9 }
```



» What about locality of instruction fetches?

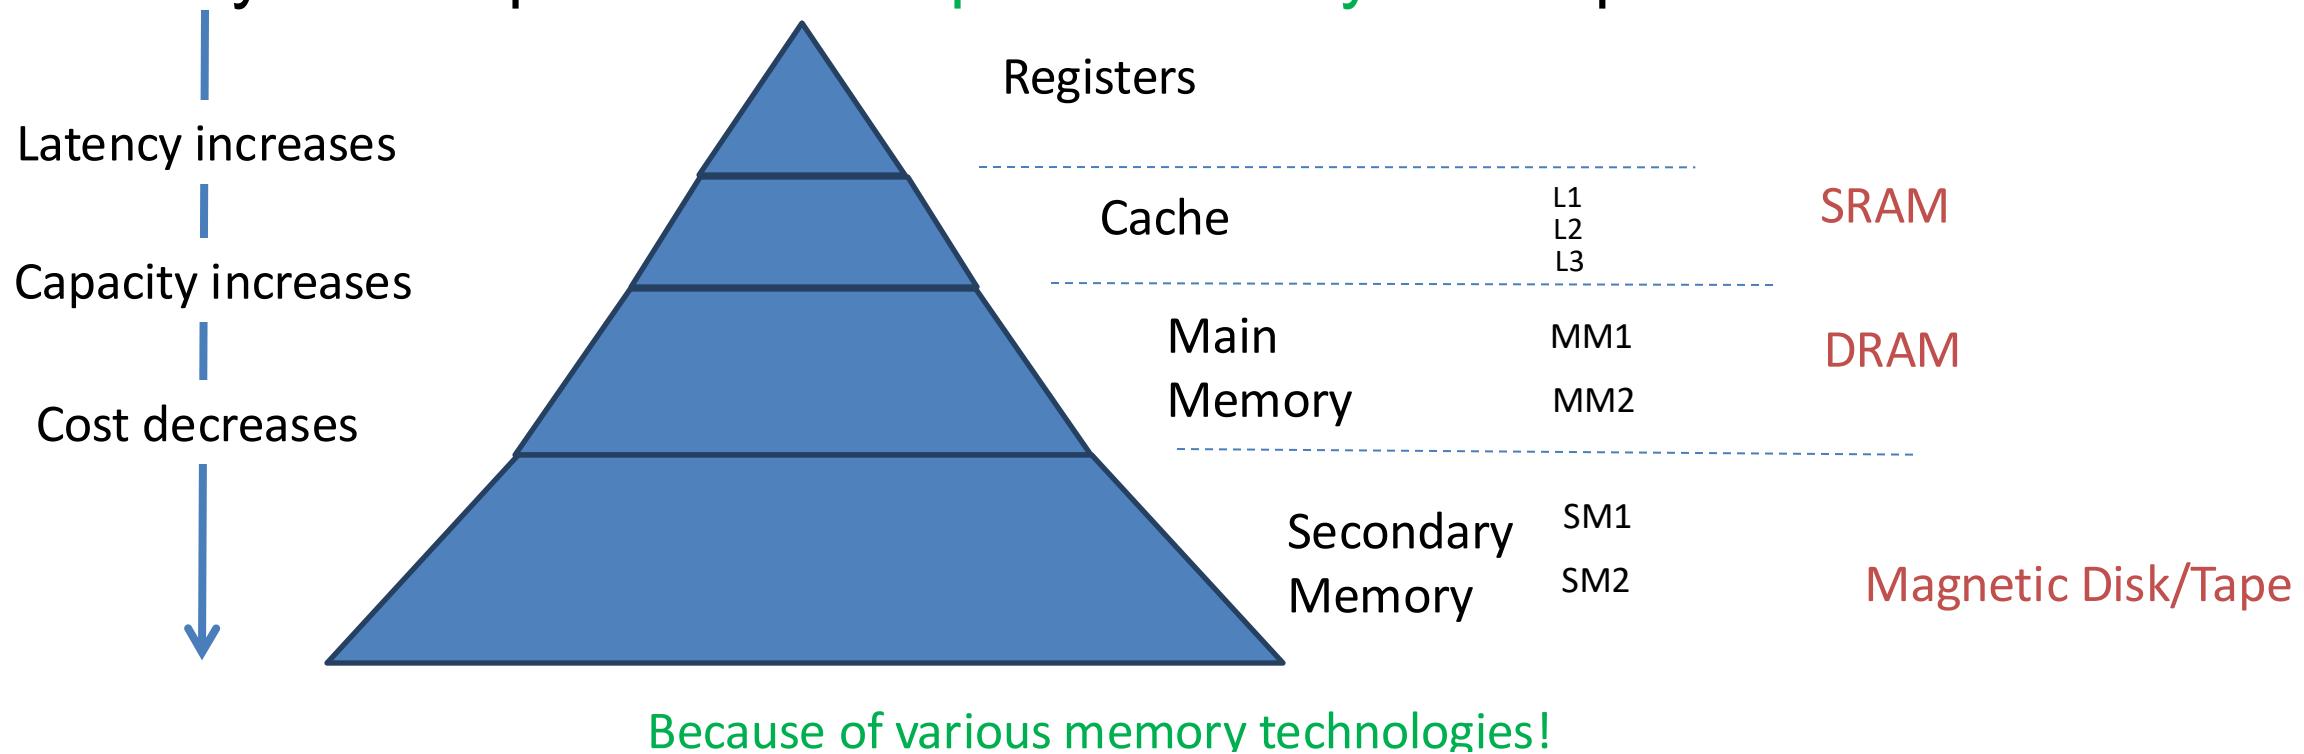
» Summary:

1. Programs that repeatedly reference the same variables exhibit good temporal locality.
2. Programs with k-stride access pattern, smaller the size better with the spatial locality.
3. Loops have good temporal and spatial locality with respect to instruction fetches.



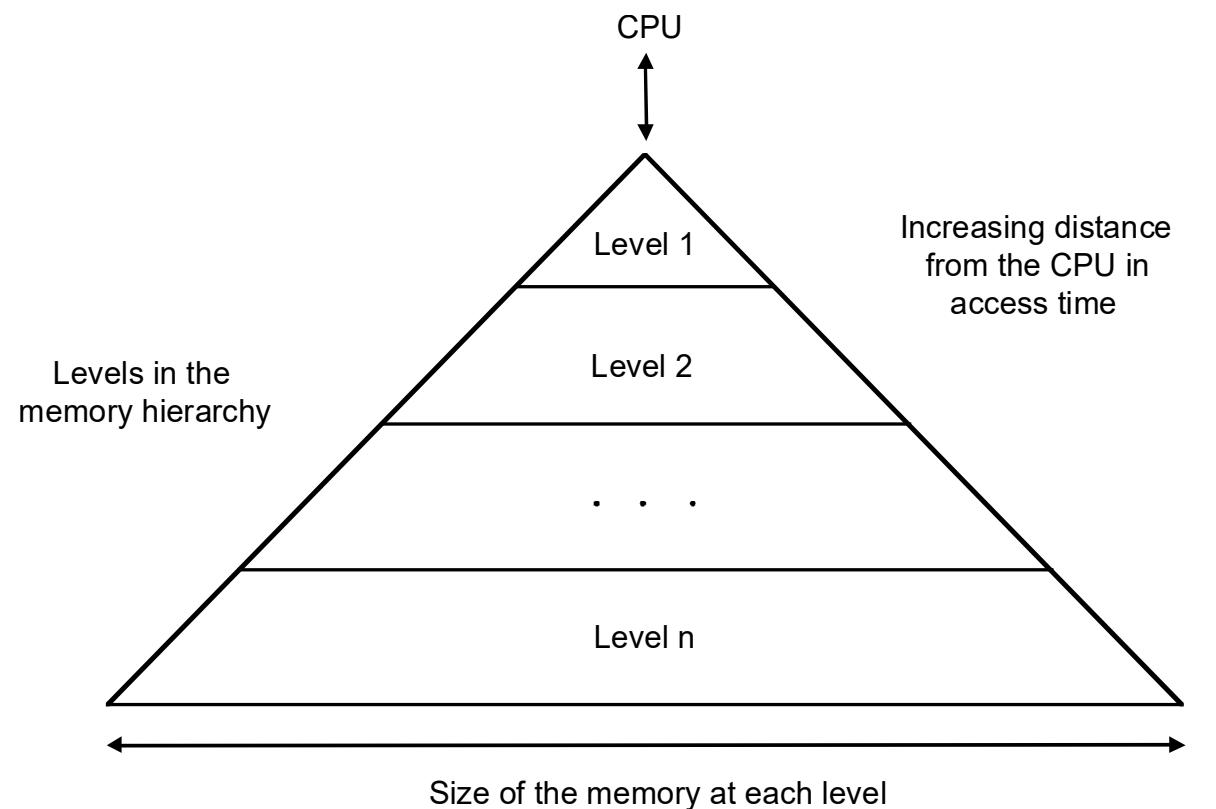
# The Memory Hierarchy!

- Practical way of achieving a fast and large memory at a low cost!
  - Heavily relies upon the '**Principle of Locality**' concept.

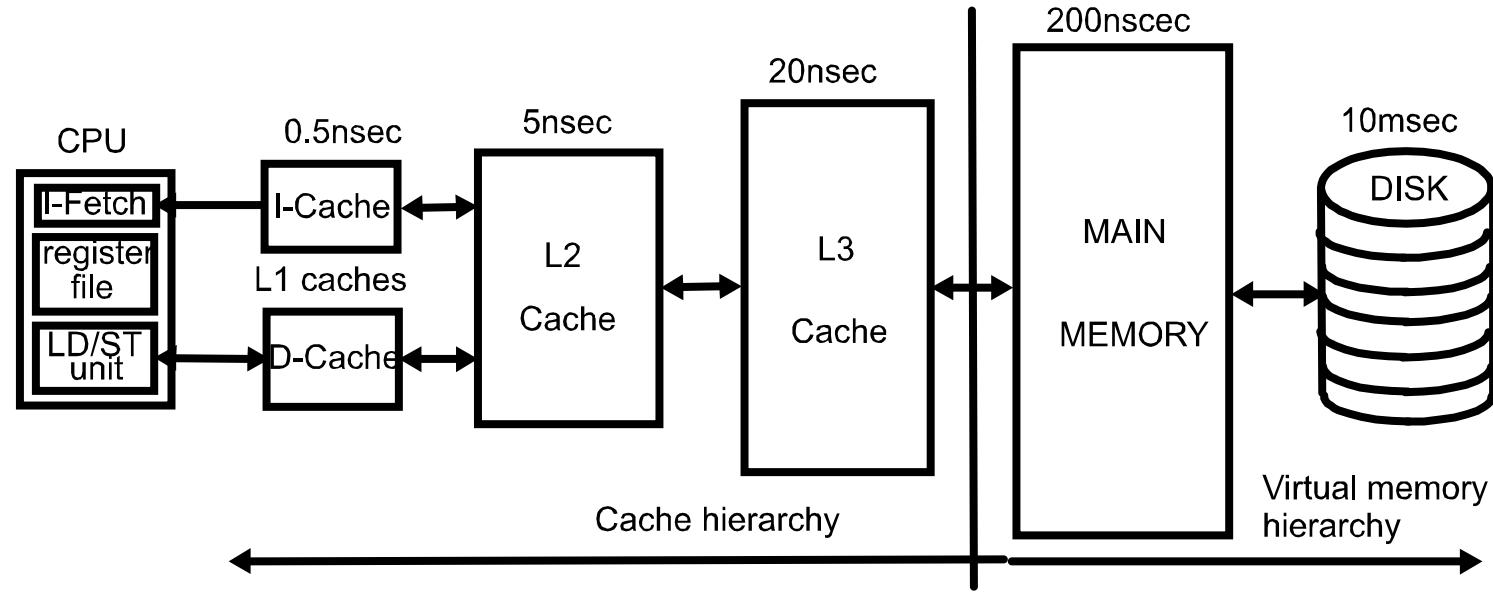


Memory technology	Typical access time	\$ per GiB in 2012
SRAM semiconductor memory	0.5–2.5 ns	\$500–\$1000
DRAM semiconductor memory	50–70 ns	\$10–\$20
Flash semiconductor memory	5,000–50,000 ns	\$0.75–\$1.00
Magnetic disk	5,000,000–20,000,000 ns	\$0.05–\$0.10

- Memory hierarchy – put small and fast memories closer to CPU, large and slow memories further away

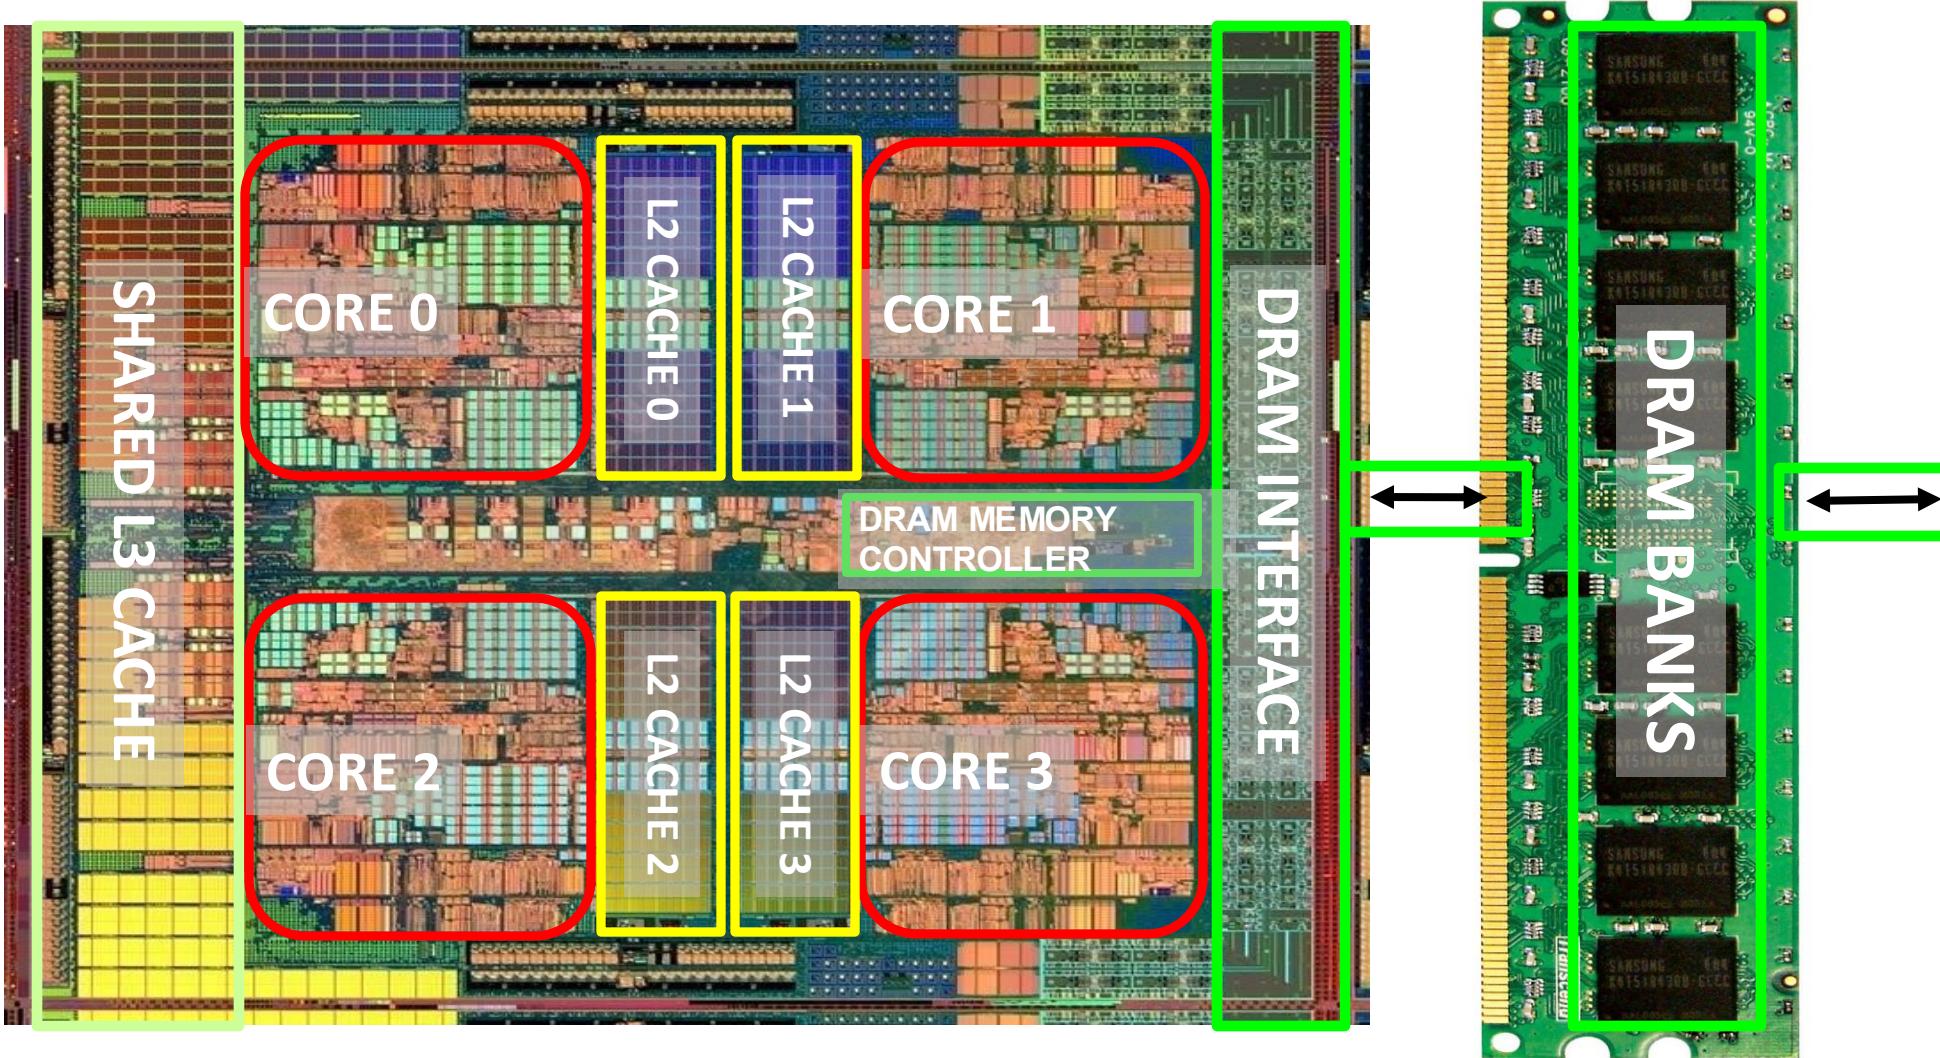


# Typical Memory Hierarchy



- Two different types of locality:
  - Temporal locality: if an item is referenced, it will tend to be referenced again soon
  - Spatial locality: if an item is referenced, items whose addresses are close tend to be referenced soon

# Typical Memory Hierarchy in a Modern System



Registers, L1 Instr & Data, L2 and L3 caches are on-chip.

# Managing Data Movement in Memory Hierarchies

## › Manual Management

- › **Programmer controls data movement** between memory levels
- ›  Useful in:
  - › Embedded processors (e.g., on-chip scratchpad SRAM instead of cache)
  - › GPUs (e.g., “shared memory”)
- ›  **Challenging for large programs**
  - › High programming effort
  - › Error-prone and hard to scale

# Managing Data Movement in Memory Hierarchies

- › **Automatic Management**
  - › Hardware handles data movement transparently
  - ›  Easier for programmers
    - › No need to understand cache size or behavior
    - › Programs remain **correct** regardless of memory hierarchy
- › **But What About Performance?**
  - › To write a **fast** program:
    - › You **must understand** the memory hierarchy
    - › Optimize data access patterns to exploit caches and memory bandwidth

# Time to pause!



# **CACHE BASICS**

# Definitions: Hits and misses

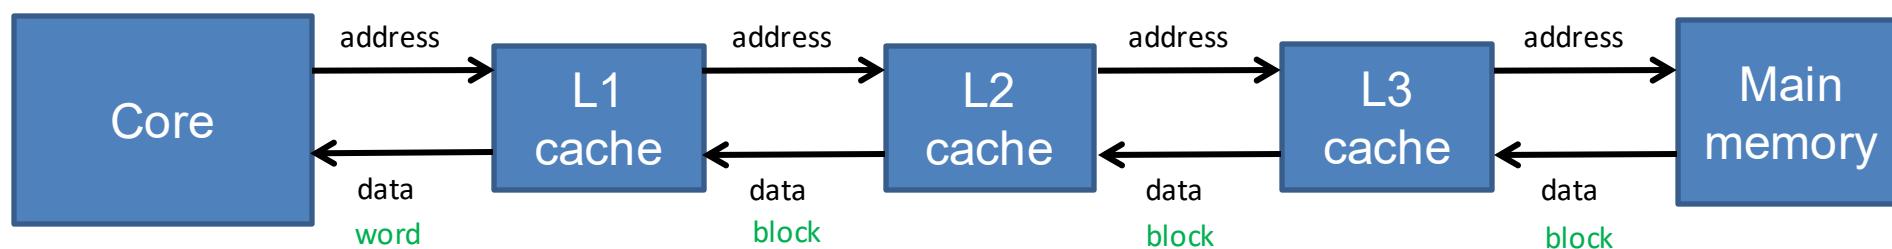
- › A **cache hit** occurs if the cache contains the data that we're looking for.
  - › Hits are good, because the cache can return the data much faster than main memory.
- › A **cache miss** occurs if the cache does not contain the requested data.
  - › This is bad, since the CPU must then wait for the slower main memory.
- › There are two basic measurements of *cache performance*.
  - › The **hit rate** is the percentage of memory accesses that are handled by the cache.
  - › The **miss rate** ( $1 - \text{hit rate}$ ) is the percentage of accesses that must be handled by the slower main RAM.

# Cache memory

- › Static RAM (SRAM) technology
  - › Each cell is made up of 6 transistors (6T)
  - › A cell stores one bit ('0' or '1')
- › 'Static' because it does not need Refresh to retain cell state.
  - › DRAM needs Refresh, we'll look at it later.
- › Volatile in nature.
- › Multiple levels of cache are most common.
  - › Working of all the levels is same.

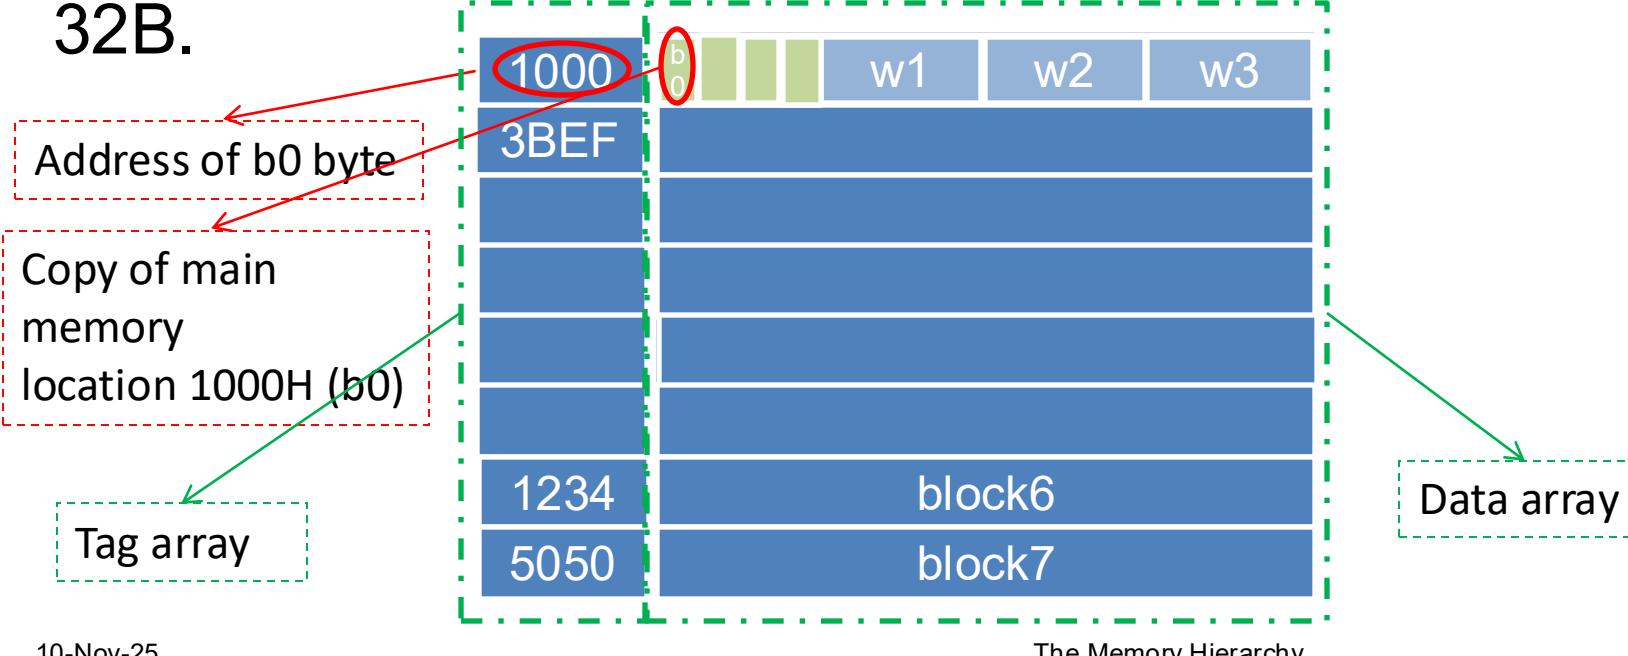
Why are we taking a block of data instead of a word?

To exploit spatial locality!!



# Cache internals

- ▶ Cache basically is a collection of blocks.
- ▶ A cache block or cache line is set of words (4B or 8B).
  - ▶ Typical cache block sizes: 32B, 64B.
- ▶ Cache size = Number of blocks \* block size
  - ▶ Ex: A 32KB cache is made up of 512 blocks of 64B each or 1024 blocks of 32B.



8 blocks of 16B.

4B words => 4 words per block.

1 word has  $4 \times 8$  SRAM cells, i.e., 32

1 block has  $4 \times 32$  SRAM cells, i.e., 128

This 8 block cache

Data array has 1024 SRAM cells.

What about Tag array?

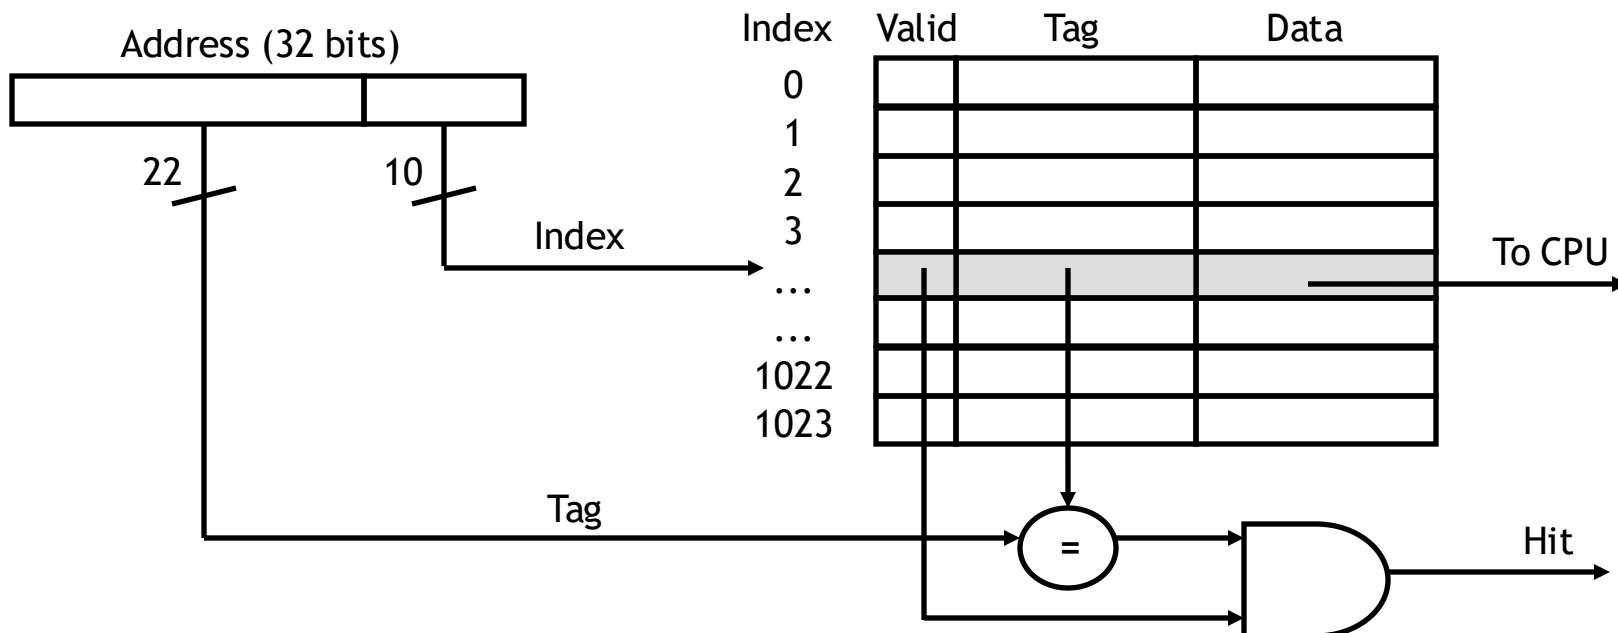
# Cache working

- Upon receiving a address from the processor, the control circuitry in the cache (namely, **cache controller**) will
  - Compare this address with those of Tag array.
    1. If there is a match, then its called **Cache Hit**.
      - Cache controller will access Data array and send a copy of the requested word to the processor.
    2. Else **Cache Miss**.
      - Cache controller will forward this address to the next level in the hierarchy.
- Four questions to understand overall working of the cache:

1. Where can a block be placed?	<b>Block placement</b>
2. How is a block found?	<b>Block identification</b>
3. Which block to replace on a miss?	<b>Block replacement</b>
4. What happens on a write?	<b>Write strategy</b>

# What happens on a cache hit

- When the CPU tries to read from memory, the address will be sent to a **cache controller**.
  - The lowest  $k$  bits of the **block** address will index a block in the cache.
  - If the block is valid and the tag matches the upper  $(m - k)$  bits of the  $m$ -bit address, then that data will be sent to the CPU.
- Here is a diagram of a 32-bit memory address and a  $2^{10}$ -byte cache.

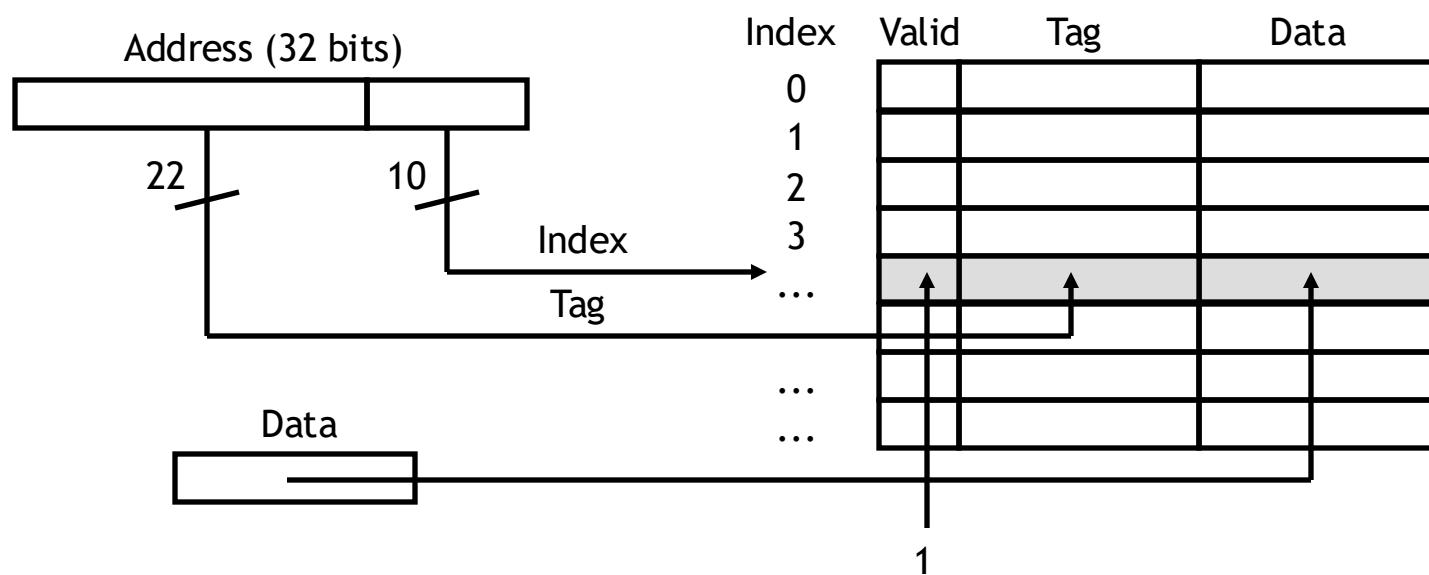


# What happens on a cache miss

- › On cache hit, CPU proceeds normally
- › On cache miss
  - › Stall the CPU pipeline
  - › Fetch block from next level of hierarchy
  - › Instruction cache miss
    - › Restart instruction fetch
  - › Data cache miss
    - › Complete data access
- › The delays that we have been assuming for memories (e.g., 1 cycle) are really assuming cache hits.

# Loading a block into the cache

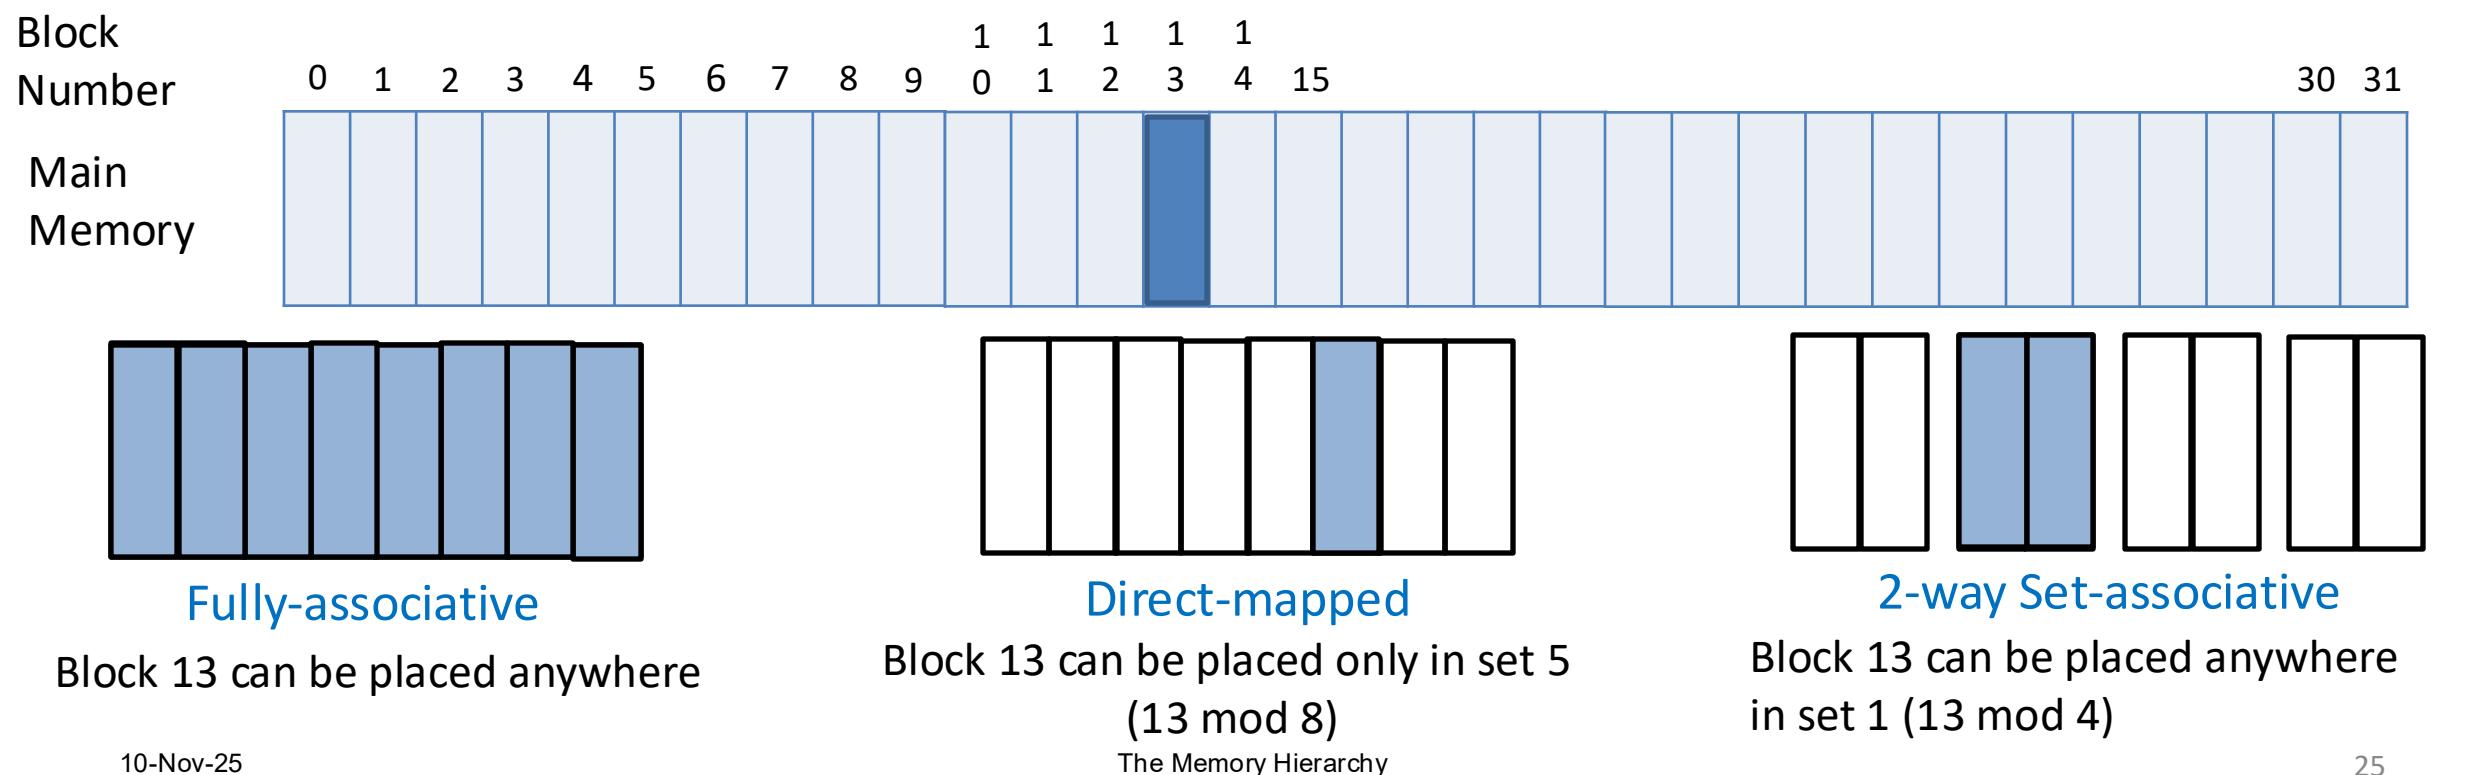
- After data is read from main memory, putting a copy of that data into the cache is straightforward.
    - The lowest  $k$  bits of the block address specify a cache block.
    - The upper  $(m - k)$  address bits are stored in the block's tag field.
    - The data from main memory is stored in the block's data field.
    - The valid bit is set to 1.



# Cache working - Q1 Block Placement

1. Where can a block be placed in cache?

1. A block can be placed anywhere in the cache Fully-associative
2. A block has only one place it can appear in cache Direct-mapped
3. A block can be placed in a restricted number of places in the cacheSet-associative
  - n possible places? n-way set associative cache.



# Cache Sets and Ways

**Ways:** Block can go anywhere

**Sets:**  
Block  
mapped  
by  
addr

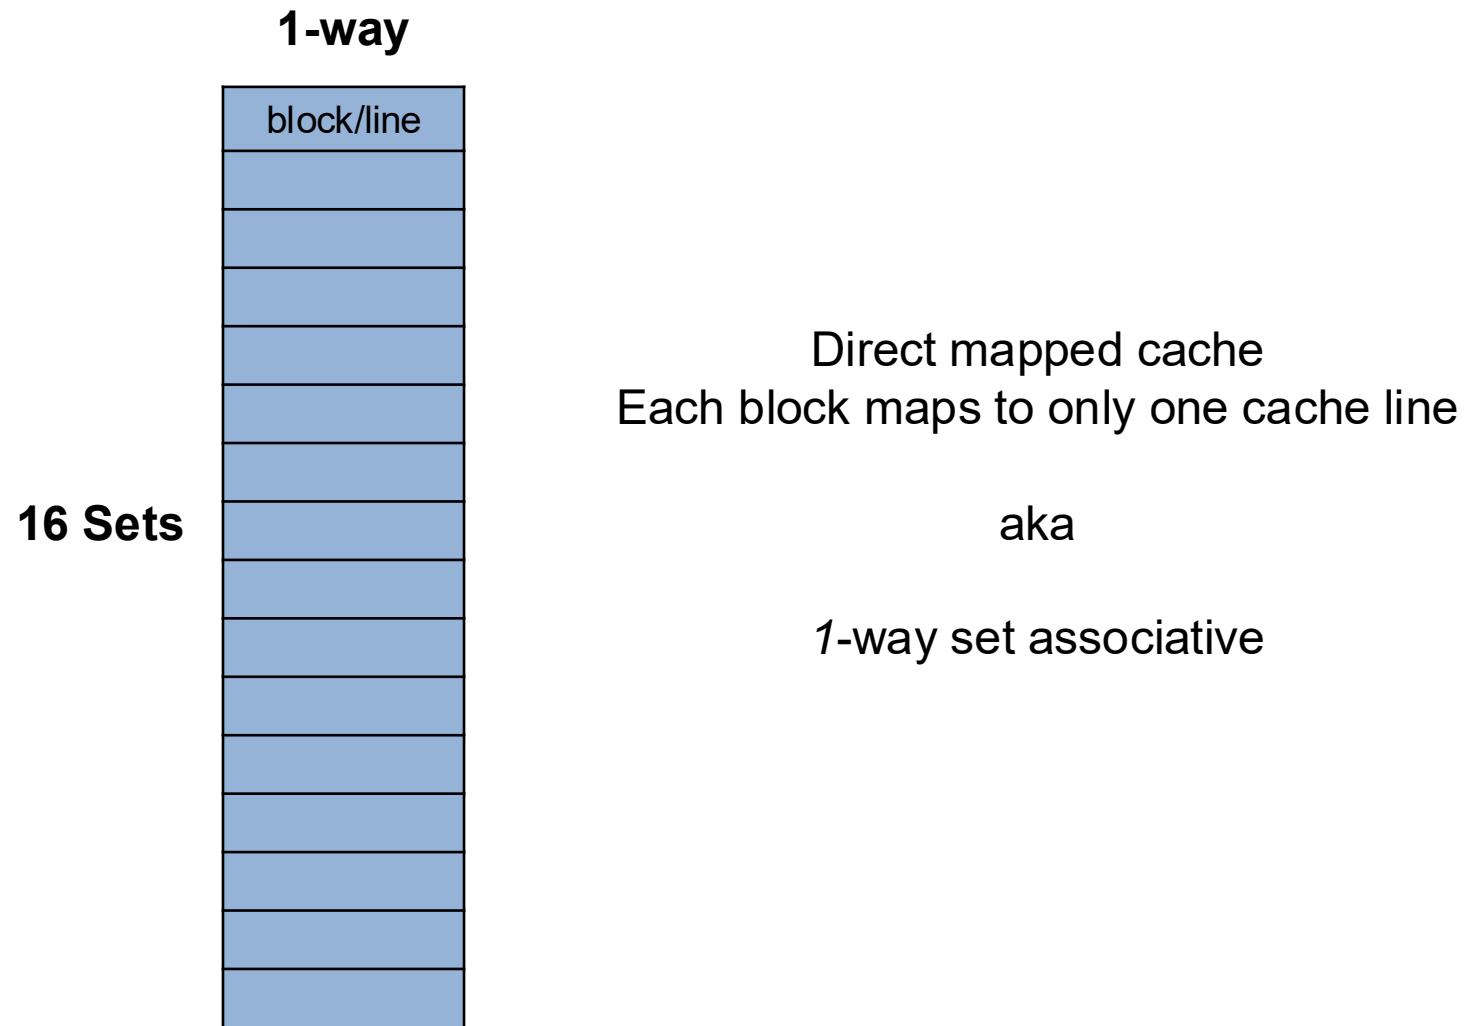
block/line			

*n*-way set associative

(4-way set associative)

Example: Cache size = 16 blocks

# Direct-mapped Cache



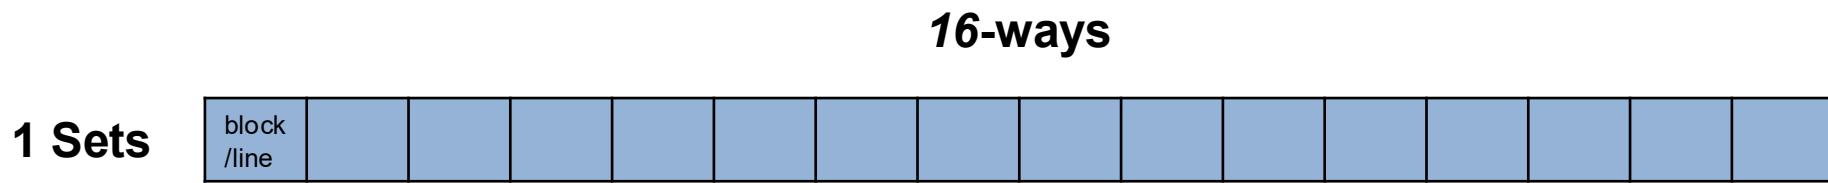
# Set Associative Cache



$n$ -way set associative  
Each block can be mapped to a set of  $n$ -lines  
Set number is based on block address

(4-way set associative)

# Fully Associative Cache

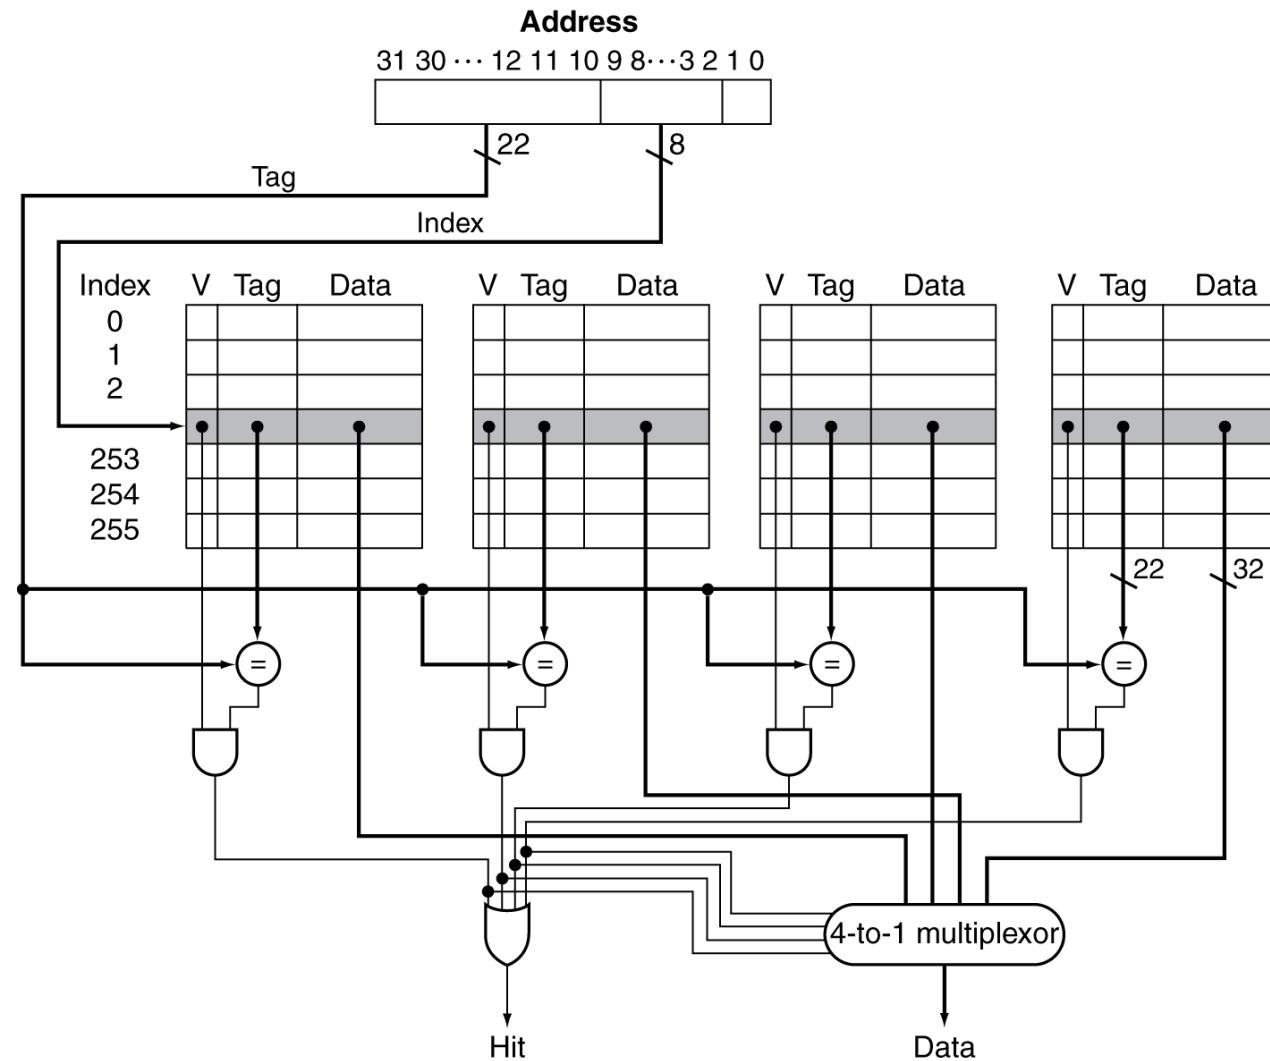


Fully associative  
Each block can be mapped to any cache line

aka

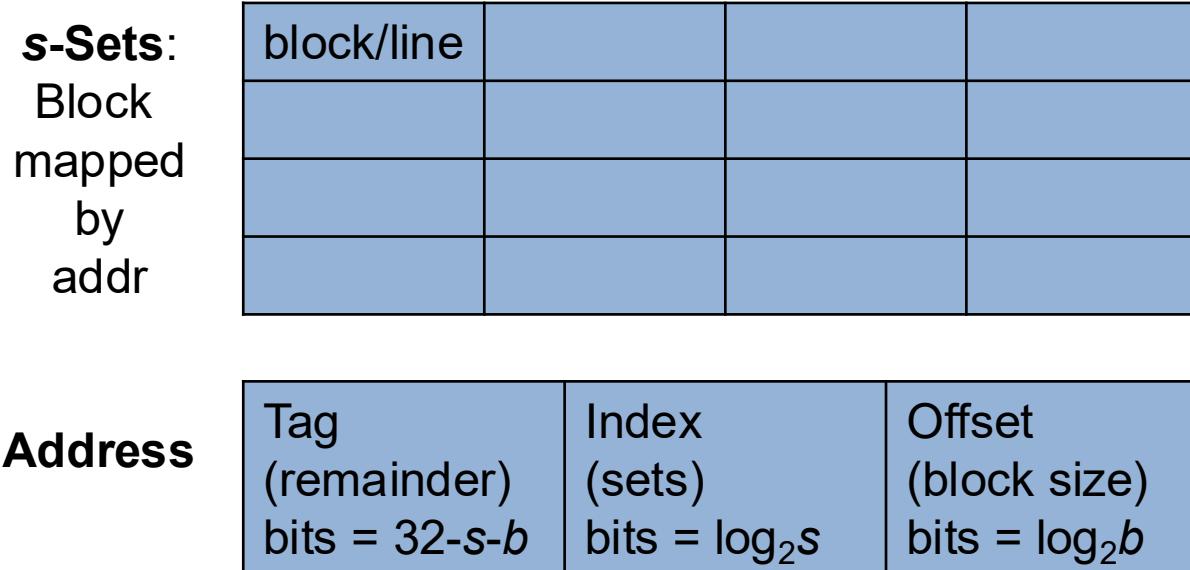
$m$ -way set associative  
where  $m$  = size of cache in blocks

# Set Associative Cache Organization



# Cache Addressing

***n-Ways***: Block can go anywhere



$s$  = number of sets

$n$  = number of ways

$b$  = block size in bytes

Cache size =  $s * n * b$   
*(area of a rectangle?)*

# Cache Addressing

Ex. 64KB cache, direct mapped, 16 byte block

Address	Tag (remainder) bits = $32-s-b$	Index (sets) bits = $\log_2 s$	Offset (block size) bits = $\log_2 b$
	16	12	4

$s$  = number of sets

$n$  = number of ways

$b$  = block size in bytes

$$\text{Cache size} = s * n * b$$

# Cache Addressing

Ex. 64KB cache, 2-way assoc., 16 byte block

Address	Tag (remainder) bits = $32-s-b$	Index (sets) bits = $\log_2 s$	Offset (block size) bits = $\log_2 b$
	17	11	4

$s$  = number of sets

$n$  = number of ways

$b$  = block size in bytes

$$\text{Cache size} = s * n * b$$

# Cache Addressing

Ex. 64KB cache, fully assoc., 16 byte block

Address	Tag (remainder) bits = $32-s-b$	Index (sets) bits = $\log_2 s$	Offset (block size) bits = $\log_2 b$
	28	0	4

$s$  = number of sets

$n$  = number of ways

$b$  = block size in bytes

$$\text{Cache size} = s * n * b$$

# What if the cache fills up?

- › What to do if we run out of space in our cache, or if we need to reuse a block for a different memory address?
- › A miss causes a new block to be loaded into the cache, automatically overwriting any previously stored data.
  - › Which block do we overwrite?
- › Cache replacement policies!

# **CACHE REPLACEMENT POLICIES**

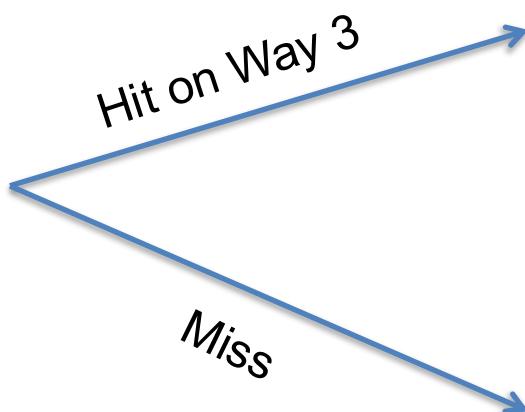
# Replacement Policy

- › Direct mapped: no choice
- › Set associative
  - › Prefer non-valid entry, if there is one
  - › Otherwise, choose among entries in the set
- › Least-recently used (LRU)
  - › Choose the one unused for the longest time
    - › Simple for 2-way, manageable for 4-way, too hard beyond that
- › Random
  - › Gives approximately the same performance as LRU for high associativity

# Cache Replacement Policies

- › Picks which block to replace within the set
- › Ex. - Random, First In First Out (FIFO), Least Recently Used (LRU), Psuedo-LRU
- › Example: LRU

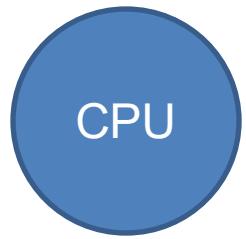
Way 0	01
Way 1	00
Way 2	11
Way 3	10



Way 0	10
Way 1	01
Way 2	11
Way 3	00

Way 0	10
Way 1	01
Way 2	00
Way 3	11

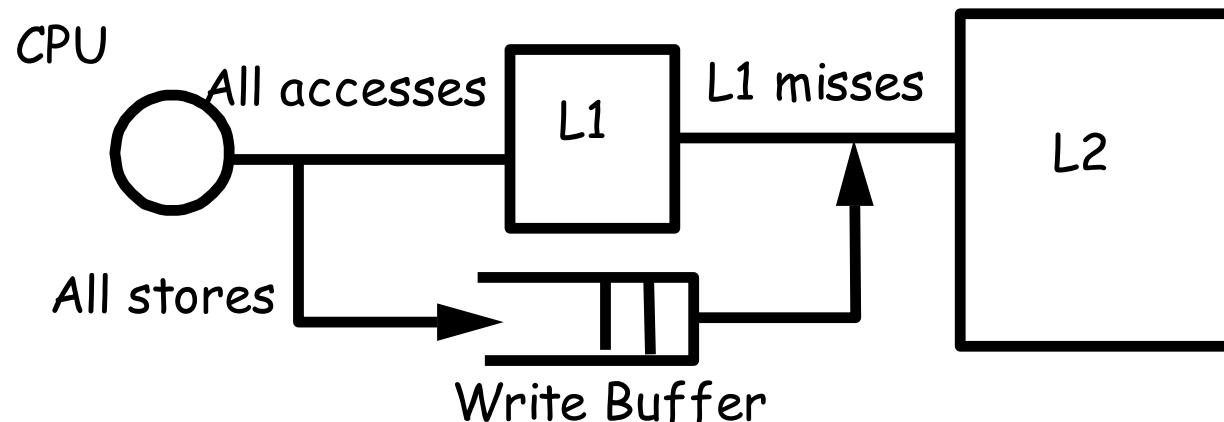
# **WRITEBACK POLICIES**



Previous slides mainly handles memory reads.  
How about memory writes?

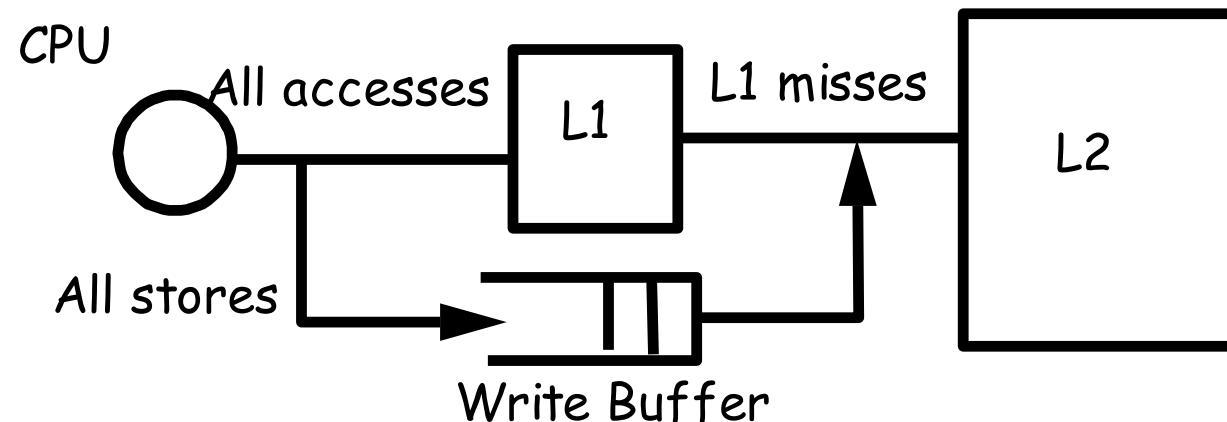
# Write-Through

- › On data-write hit, could just update the block in cache
  - › But then cache and memory would be inconsistent
- › Write through: also update memory



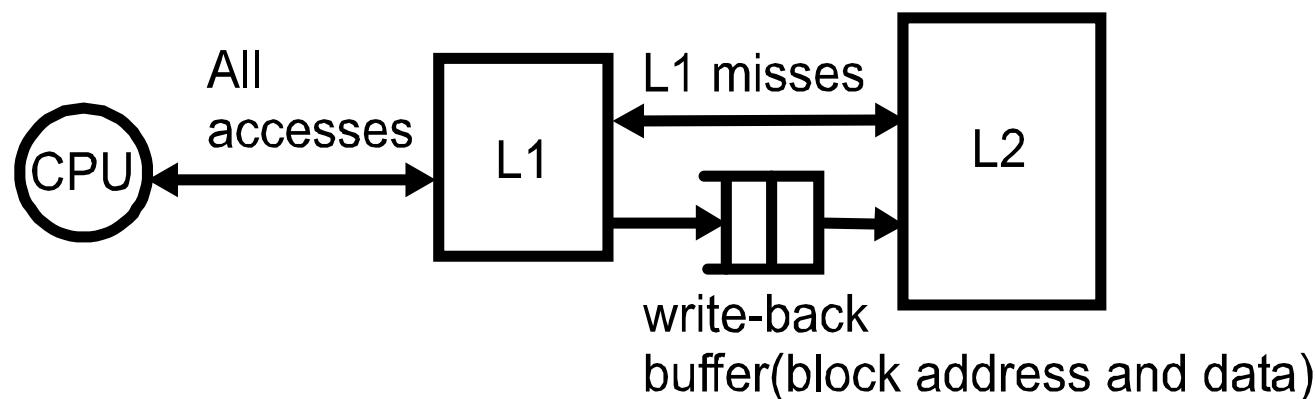
# Write-Through

- › But makes writes take longer
  - › e.g., if base CPI = 1, 10% of instructions are stores, write to memory takes 100 cycles
    - › Average CPI =  $1 + 0.1 \times 100 = 11$
- › Solution: write buffer (Holds data waiting to be written to memory)
  - › CPU continues immediately
    - › Only stalls on write if write buffer is already full



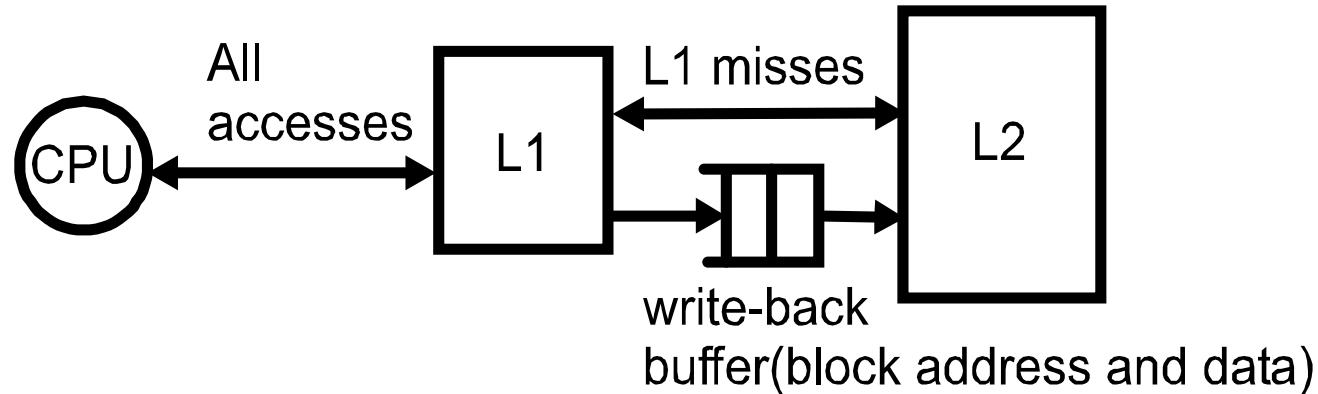
# Write-Back

- › Alternative: On data-write hit, just update the block in cache
  - › Keep track of whether each block is dirty
- › When a dirty block is replaced
  - › Write it back to memory
  - › Can use a write buffer to allow replacing block to be read first



# Write Allocation

- › What should happen on a write miss?



- › For write-back: Usually fetch the block
- › Alternatives for write-through
  - › Allocate on miss: fetch the block
  - › Write around: don't fetch the block

# **CACHE PERFORMANCE**

# Measuring Cache Performance

- › Components of CPU time
  - › Program execution cycles: Includes cache hit time
  - › Memory stall cycles: Mainly from cache misses
- › With simplifying assumptions:

$$\begin{aligned}\text{Memory stall cycles} &= \frac{\text{Memory accesses}}{\text{Program}} \square \text{Miss rate} \square \text{Miss penalty} \\ &= \frac{\text{Instructions}}{\text{Program}} \square \frac{\text{Misses}}{\text{Instruction}} \square \text{Miss penalty}\end{aligned}$$

# Measuring Cache Performance

- › I-cache miss rate = 2%, D-cache miss rate = 4%, Miss penalty = 100 cycles,  
Base CPI (ideal cache) = 2,  
Load & stores are 36% of instructions
- › What is the actual CPI?
- › Memory stall cycles per instruction to:
  - › D-cache:  $0.36 \times 0.04 \times 100 = 1.44$
  - › I-cache:  $1.0 \times 0.02 \times 100 = 2$
- › Actual CPI = 2 + 2 + 1.44 = 5.44

$$\text{Memory stall cycles} = \frac{\text{Memory accesses}}{\text{Program}} \square \text{Miss rate} \square \text{Miss penalty}$$
$$= \frac{\text{Instructions}}{\text{Program}} \square \frac{\text{Misses}}{\text{Instruction}} \square \text{Miss penalty}$$

# Average Access Time

- › Hit time is also important for performance
- › Average memory access time (AMAT)
  - ›  $AMAT = \text{Hit time} + \text{Miss rate} \times \text{Miss penalty}$
- › Example
  - › CPU with 1ns clock, hit time = 1 cycle,  
miss penalty = 20 cycles, I-cache miss rate = 5%
  - ›  $AMAT = 1 + 0.05 \times 20 = 2\text{ns}$



# Progress so far .....

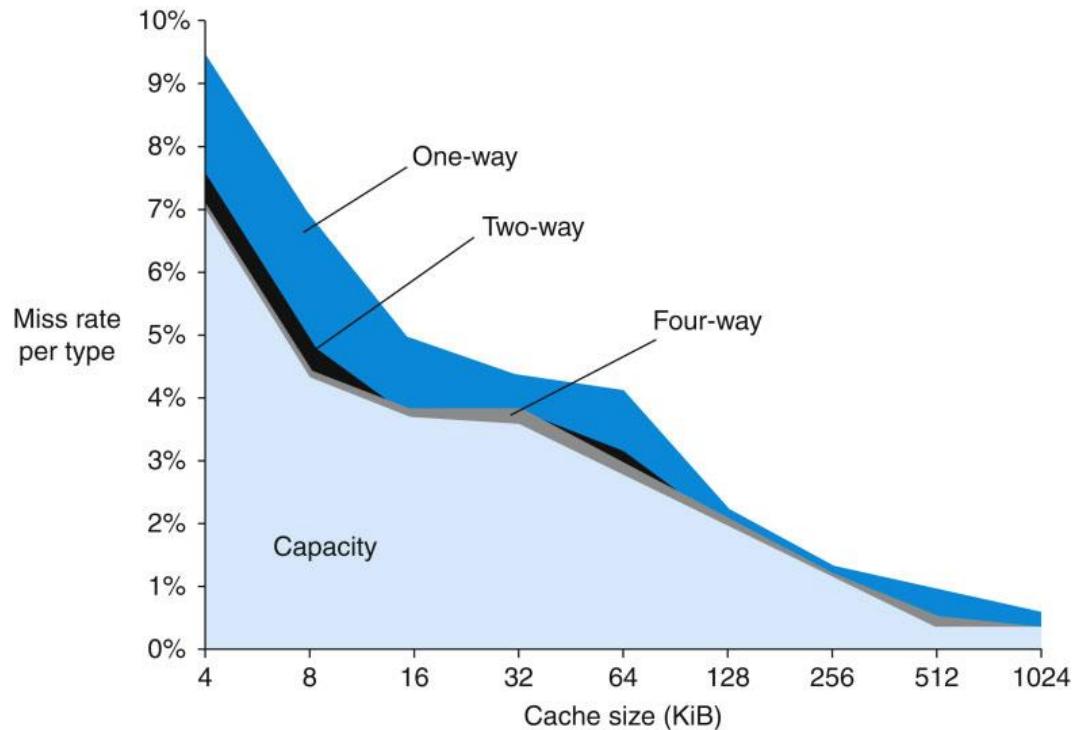
- Four questions to understand overall working of the cache:
  1. Where can a block be placed? **Block Placement**
    - Cache organization and Design
    - Cache Addressing , sets, ways
  2. How is a block found? **Block Identification**
    - Measure Cache Performance, Average Memory Access Time (AMAT)
  3. Which block to replace on a miss? **Replacement Policies**
    - Depends on Cache Organization [Direct/Set Associative/Fully Associative]
    - FIFO, LRU (Recently), Random, LFU (Frequently)
  4. What happens on a write? **Write Back Policies**
    - Write back/ Write Through
    - Write Allocation

# Cache Performance

- › Miss rate
  - › Fraction of cache access that result in a miss
- › Causes of misses (3C's)
  - › Compulsory (cold)
    - › First reference to a block
  - › Capacity
    - › Space is not sufficient to host data or code
  - › Conflict
    - › When two memory blocks map on the same cache block in direct-mapped or set-associative caches

# Sources of Misses – 3C's

- › Compulsory misses (aka cold start misses)
  - › **First access** to a block
  - › Unavoidable when data is accessed for the first time
- › Capacity misses
  - › Due to **finite cache size**
  - › A **replaced block** is later accessed again
- › Conflict misses (aka collision misses)
  - › In a **non-fully associative cache**
  - › Due to competition for entries in a set
  - › **Would not occur** in a fully associative cache of the same total size



# Measuring/Classifying Misses

- › How to find out each type?
- › Cold misses:
  - › Simulate a fully-associative infinite cache size
- › Capacity misses:
  - › Simulate fully-associative cache, then deduct cold misses
- › Conflict misses:
  - › Simulate target cache configuration then deduct cold and capacity misses

# Measuring/Classifying Misses

- › Classification is useful to understand how to eliminate misses
- › High capacity misses
  - › → need larger cache
- › High conflict misses
  - › → need higher associativity

# **CACHE OPTIMIZATIONS**

# Think of Cache Like a Fast-Access Bookshelf



Small and simple first-level caches

- › Like keeping only your most-used books on a small shelf near your desk for quick access

Critical timing path = Steps to grab the right book

1. Find the shelf (addressing tag memory)
2. Check the book title (comparing tags)
3. Pick-the correct book (selecting the right set)

Direct-mapped caches

- › Like having only one spot for each book—easy to check and grab at the same time

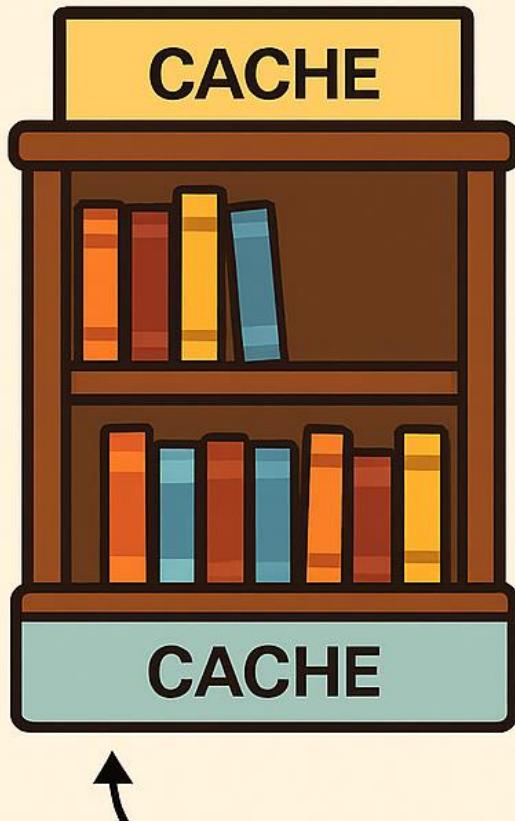
Lower associativity saves power

- › Fewer shelves to scan means less effort (fewer cache lines accessed)

# Measuring/Classifying Misses

- › Classification is useful to understand how to eliminate misses
- › High capacity misses
  - › → need larger cache

# Reduce Capacity Misses: Make the Shelf Bigger



## Trade-offs

- Longer hit time: Bigger shelf = more scanning before you find the book

Idea: Increase the cache size so more data fits inside

Imagine your desk bookshelf is too small, and you keep running to the main library. Adding a bigger shelf means fewer trips—but it takes more space, costs more, and might slow you down when searching

## Trade-offs

- Longer hit time: Bigger shelf = more scanning before you find the book
- Higher cost and power: Larger furniture needs more material and energy
- Popular for off-chip caches: Like adding a big storage cabinet outside your room—great for bulk storage



For off-chip caches

# Objective: Minimize Miss Rate and Miss Penalty

Think of your cache like a personal bookshelf. The goal is to avoid running to the library (main memory) and save time when you do.

## 1. Increase Block Size

→ Bring Bigger Book Bundles

✓ Reduces compulsory misses

⚠ Like borrowing a whole series instead of one book – you need fewer trips

Drawbacks: Bigger bundles take more space and time to carry → may increase capacity & conflict misses and miss penalty



## 3. Increase Associativity

→ Add More Slots per Shelf

✓ Reduces conflict misses

⚠ Like having flexible spots for books – less chance of eviction

Drawbacks: More complexity = slower lookup and higher energy → increases hit time and power consumption

## 3. Increase Associativity

→ Add More Slots per Shelf

⚠ Like having flexible spots for books – less chance of

# Cache Optimization Objective: Minimize the miss rate

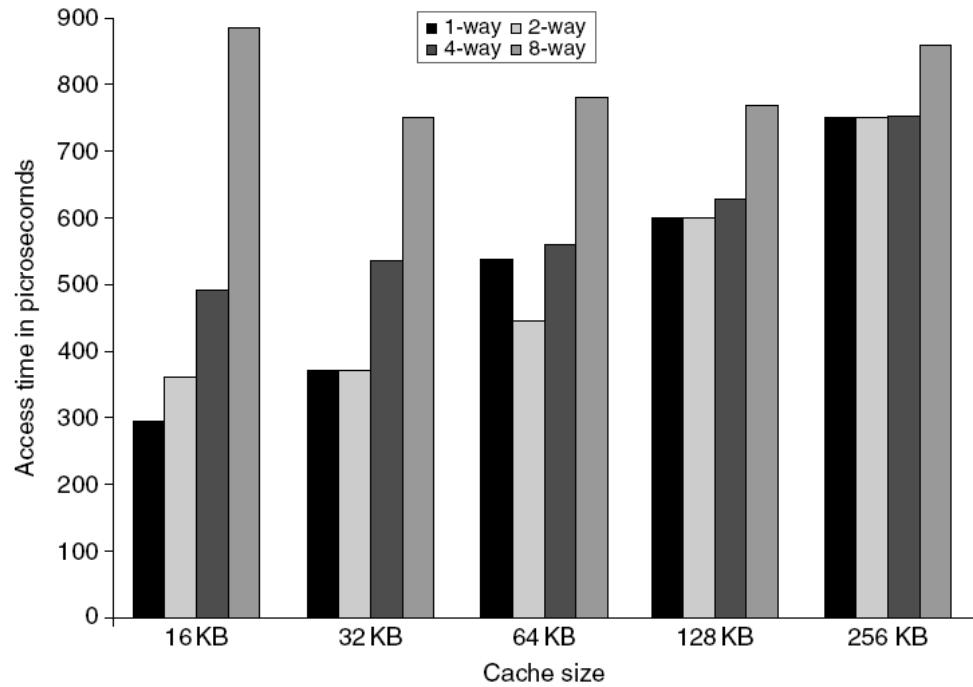
- **1. Increase Block Size**
  -  Reduces **compulsory misses**,  May increase **capacity & conflict misses**, and **miss penalty**
  - **Block Size** → **Bigger book bundles** (fewer trips, but heavier and slower).
- **2. Increase Total Cache Capacity**
  -  Lowers **miss rate**,  Raises **hit time** and **power consumption**
  - **Cache Capacity** → **Bigger shelf** (more books, but harder to search and costly).
- **3. Increase Associativity**
  -  Reduces **conflict misses**,  Increases **hit time** and **power consumption**
  - **Associativity** → **Flexible slots** (reduces conflicts, but adds complexity and power cost).

# Cache Optimizations Basics

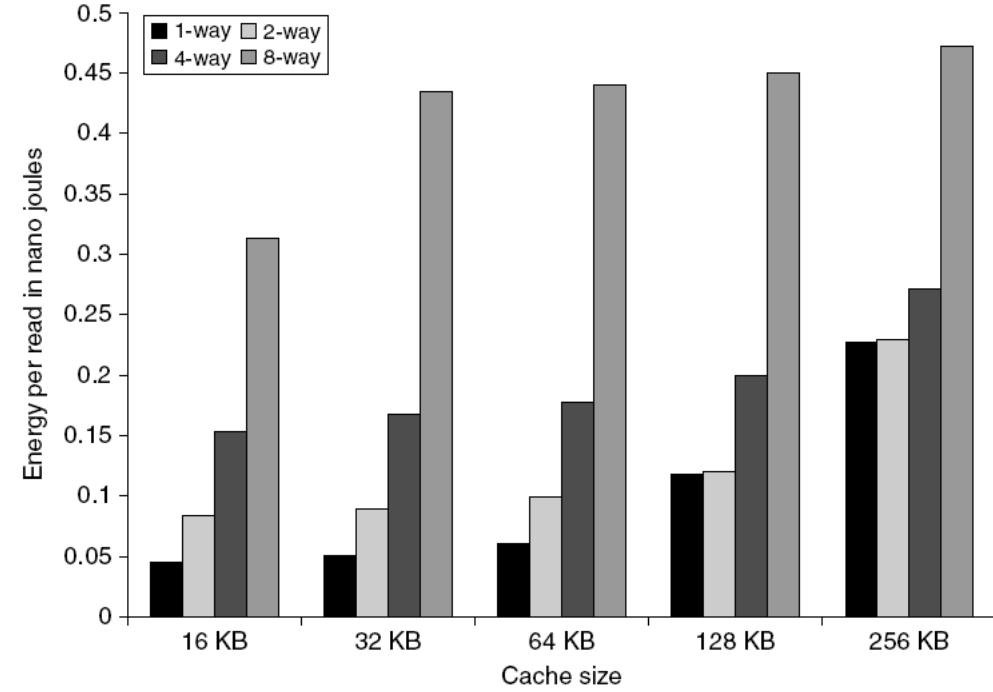
Objective: Minimize the miss rate and miss penalty

- › **4. Add More Cache Levels**
  - ›  Reduces overall memory access time
- › **5. Prioritize Read Misses Over Writes**
  - ›  Lowers miss penalty
- › **6. Avoid Address Translation in Cache Indexing**
  - ›  Reduces hit time

# L1 Size and Associativity



Access time vs. size and associativity



Energy per read vs. size and associativity

# Optimization 4: Multi-Level Caches

Focus: Miss Penalty

# 4. Add More Cache Levels → Create a Chain of Shelves

- ›  **Reduces overall memory access time**
- › **Analogy:**
  - › Imagine you have a small shelf on your desk (L1),
  - › a bigger shelf in your room (L2),
  - › and a large cabinet in the hallway (L3).
  - › Instead of running all the way to the library (main memory), you check the closest shelf first, then the next, and so on.
  - › **Result:** Faster access because you rarely need to go to the library.

# What Should be Our Focus?

- › **Reduce Miss Penalty OR Reduce Miss Rate**
  - › **Memory Wall Problem:**  
Processors are sprinting ahead, but memory is lagging behind.
- Question:**
- › Should we **make the cache faster** (keep pace with the runner)?
  - › Or **make the cache bigger** (build more rest stops along the track)?

# Inclusive Cache → “Everything in L1 is also in L2”

- › **Analogy:**  
Think of L1 as your desk shelf and L2 as a room bookshelf.  
Every book on your desk is also in the room shelf.
- › **Why desirable?**  
Easy to check consistency—just look at the bigger shelf (L2).
- › **Block size issue:**  
If desk shelf uses small books (64B) and room shelf uses big books (128B), removing one big book means clearing multiple small ones → messy!
- › **Solution:**  
Keep block sizes the same for both shelves.

# My Chip space can Accommodate only a **SMALL L2 cache**

- › **Exclusive Cache** → “**No duplicates between shelves**”
- › **Analogy:**
  - › Desk shelf and room shelf never hold the same book. If you miss on the desk shelf, swap a book from the room shelf instead of adding a duplicate.
- › **Benefit:**
  - › Saves space on the big shelf.
- › **Example:**
  - › AMD Opteron uses two 64 KiB L1 shelves and a 1 MiB L2 shelf with this policy.

# Optimization 5: Read Misses prioritized over Writes

## Focus: Miss Penalty

# 5. Prioritize Read Misses Over Writes → Serve Readers First

- ›  **Lowers miss penalty**
- › **Analogy:**
  - › Think of a bookstore where customers are waiting. People who want to **read now** (read misses) get priority over those who want to **return books or make changes** (writes).
  - › Why? Reading is urgent for progress, while writing can wait a little without hurting performance.

# 5. Prioritize Read Misses Over Writes → Serve Readers First

- ›  **Lowers miss penalty**
- › **Analogy:**
  - › Think of a bookstore where customers are waiting. People who want to **read now** (read misses) get priority over those who want to **return books or make changes** (writes).
  - › Why? Reading is urgent for progress, while writing can wait a little without hurting performance.

# 5. Prioritize Read Misses Over Writes

- Serve reads before writes have been completed.
  - **Analogy:**
    - Imagine a busy café: customers who want to **read the menu and order** (read requests) are served first, even if the chef is still finishing previous dishes (writes). This keeps the line moving.
- Write buffer complexity:
  - **Question:** How big should the buffer be in a write-through cache?
  - **Purpose:** Hold updated values that might be needed during a read miss

# 5. Prioritize Read Misses Over Writes

## ➤ Two Approaches

### ➤ Wait for the buffer to empty

- Like waiting for the chef to finish all pending dishes before taking a new order—slow and inefficient.
- **Read miss waits until the write buffer is empty**

### ➤ Check the buffer and proceed if safe

- If no conflicts and the kitchen is ready, take the new order immediately.
- **Check the contents of the write buffer on a read miss, and if there are no conflicts and the memory system is available, let the read miss continue**
-  **All modern systems use this approach** for speed.

Optimization 6: Avoid Address Translation  
during Indexing of the Cache to Reduce Hit  
Time.

Focus: Hit Time

# Address Translation During Cache Indexing

- Refers to the step where the processor converts a **virtual address** (used by programs) into a **physical address** (used by memory hardware) before using it to index into the cache.
- **Why does this matter?**
  - Most modern systems use **virtual memory**, so the CPU generates virtual addresses.
  - The cache typically works with **physical addresses** for correctness (especially in multiprogramming).
  - If you wait for the **translation** (via the TLB – Translation Lookaside Buffer) before indexing the cache, it adds extra delay to the **hit time**.
    - Virtual address → TLB → Physical address → Cache indexing (slow path)
    - Optimized path: Virtual index first, physical tag later?

# Address Translation During Cache Indexing

- **Analogy:**
  - Imagine you have a **nickname** (virtual address) and a **legal name** (physical address).
  - If the librarian insists on checking your legal name before letting you grab a book, you waste time.
- **Optimization:**
  - Use the nickname directly to find the shelf (index the cache) while verifying later if it matches the legal name (tag check).
  - This technique is called **Virtually Indexed, Physically Tagged (VIPT)** caching:
    - **Index with virtual address** (fast).
    - **Tag compare with physical address** (ensures correctness).

# Topics Covered..... UG Syllabus

- Four questions to understand overall working of the cache:
  1. Where can a block be placed? **Block Placement**
    - Cache organization and Design
    - Cache Addressing , sets, ways
  2. How is a block found? **Block Identification**
    - Measure Cache Performance, Average Memory Access Time (AMAT)
  3. Which block to replace on a miss? **Replacement Policies**
    - Depends on Cache Organization [Direct/Set Associative/Fully Associative]
    - FIFO, LRU (Recently), Random, LFU (Frequently)
  4. What happens on a write? **Write Back Policies**
    - Write back/ Write Through
    - Write Allocation

# Topics Covered..... PG Syllabus

- › Different types of cache misses and how to measure them
- › Metrics in Cache Performance
  - ›  $AMAT = \text{Hit time} + \text{Miss rate} \times \text{Miss penalty}$
- › Cache optimizations
  - › **Miss Rate:** Increase Block Size/Total Cache Capacity/Associativity
  - › **Miss Penalty:** More Cache Levels, Prioritize Read Misses Over Writes
  - › **Hit Time:** Avoid Address Translation in Cache Indexing

# **End of Cache Topic/Syllabus**