

# Challenges and opportunities in ML accelerator architectures

Kiran Ranganath  
*SambaNova Systems*

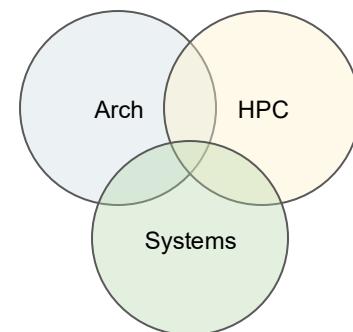
# Bio

Principal Researcher, SambaNova Systems

PhD, UC Riverside

Advisor: Dr. Daniel Wong

Research interests: Computer architecture, high-performance computing, Systems for ML



# Introduction

# Dawn of transformer-based LLMs

- OpenAI - GPT, GPT 2 (1.5B), GPT 3 (13B), GPT 4 etc.
- Meta - Llama, Llama 2 (7B, 13B, 70B), Llama 3.x (8B, 70B, 405B), Llama 4 (100B, 400B)
- Mistral - 3B, 7B, 8B, 8x22B etc.
- BigScience - Bloom (176B)
- DeepSeek (671B)
- *Many more..*

We need the ability to **train** and **serve** models at scale.

Train = Pre-train + Fine-tune

Serve = Inference

# Transformers for NLP

[NeurIPS'17] Attention is all you need<sup>[1]</sup>

Key takeaways for HPCA enthusiasts:

- **Fast** to train
- Support for **longer sequence** lengths/contexts
- It converted NLP into a **Highly parallelizable** compute problem

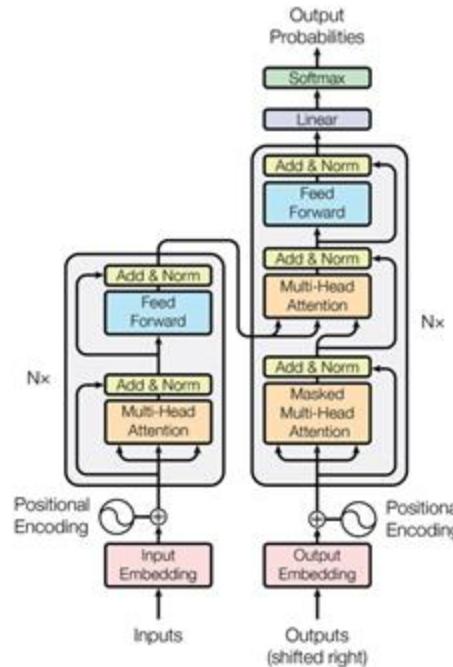


Figure 1: The Transformer - model architecture.

# Breaking down the transformer model architecture

- $N \times$  decoders
- Each decoder
  - Multi-head attention (MHA) ( $h \times$  heads)
    - Scaled Dot Product Attention (SDPA)
    - Q, K, V Linear
- Complexity?
  - Bulk :  $N \times h \times (SDPA + Q + K + V + \text{Feed forward})$
  - GeMM (Linear & MatMul) throughput dominates model performance.

DeepSeek v3

$N = 61$

$h = 128$

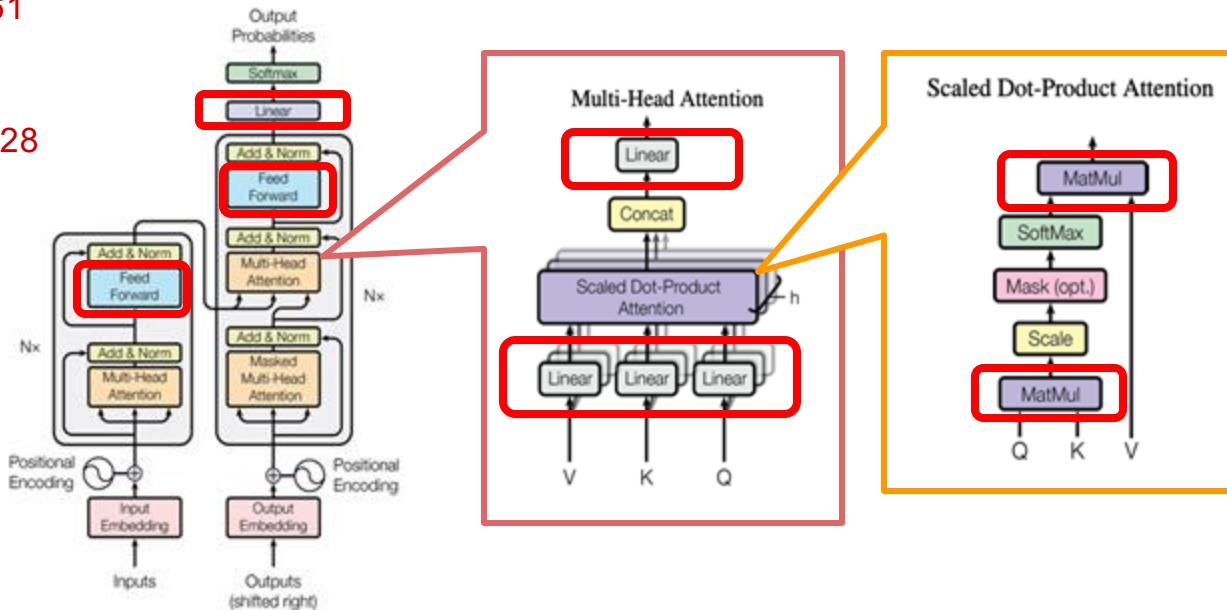


Figure 1: The Transformer - model architecture.

# Challenge 1: ML as an HPC problem

# Speeding up General Matrix Multiplication (GeMMs)

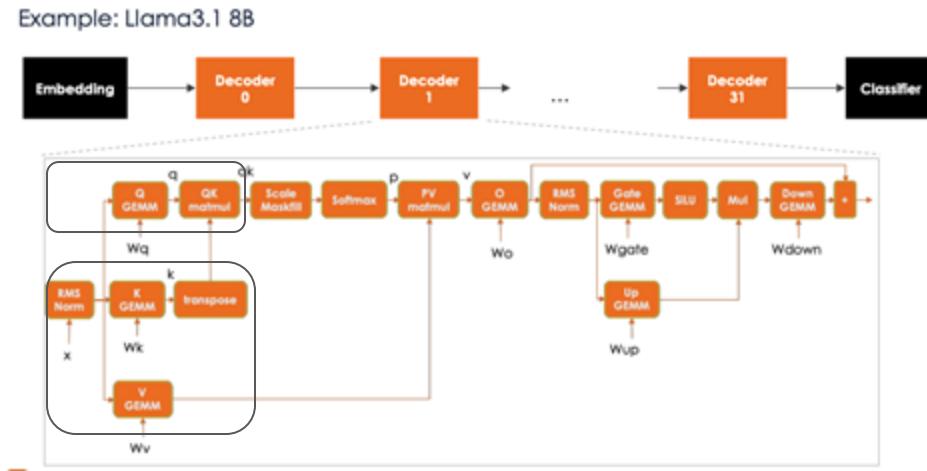
- CPU optimizations?
  - Effective use of caches, locality, etc.
  - Libraries - OpenBLAS etc.
  - Not suitable for large matrices
  - Can extract throughputs of <1 TFLOPs per chip
- GPU optimizations?
  - Effective use of global memory, shared memory, high-bandwidth memory (HBM), tiling, etc.
  - Proprietary libraries like CuBLAS offer near to linear scaling.
  - GeMMs can be scaled across multiple GPUs
  - Throughputs can range from 10+ TFLOPs per chip

```
#include <stdio.h>

// Function to perform GeMM: C = alpha * A * B + beta * C
void gemm(int M, int N, int K,
          float alpha, float *A, float *B,
          float beta, float *C)
{
    for (int i = 0; i < M; ++i) {
        for (int j = 0; j < N; ++j) {
            float sum = 0.0f;
            for (int k = 0; k < K; ++k)
                sum += A[i*K + k] * B[k*N + j];
        }
        C[i*N + j] = alpha * sum + beta * C[i*N + j];
    }
}
```

# Breaking down Llama 3.1 8B inference

- 32 decoders
- 9 GeMMs per encoder
- Can support 128k sequence length
- Model parameters = ~16 GB (BF16)
- KV cache size = up to 4GB / 1k sequence / batch
- For sequence length = 16k,  
**Memory required = ~1TB**  
**Model FLOPs = ~8.8 TFLOPs**



# Need of the hour

- **Faster & larger on-device memory**
  - High-bandwidth memory (HBM), SRAM etc.
- **Higher TFLOPs**
  - Higher core frequency
  - Pack more compute threads

# Challenge 2: Thinking beyond GPUs

# Breaking down Llama 3.1 8B inference

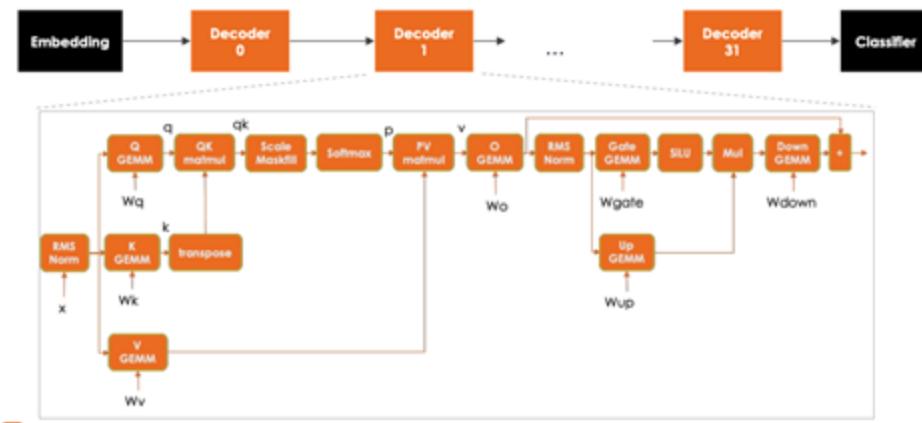
Llama 3.1 8B config:

**Memory required = 1TB**  
**Model FLOPs = ~8.8 TFLOPs**

How to execute this graph?

- Block by block?
- Section by section?

Example: Llama3.1 8B

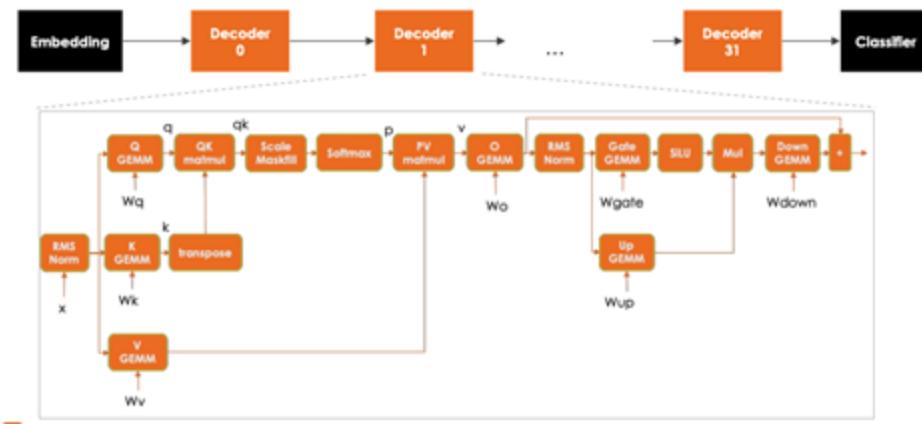


# Block by block

Execute each of the blocks  
individually one-by-one  
(kernel-by-kernel) - GPU-style

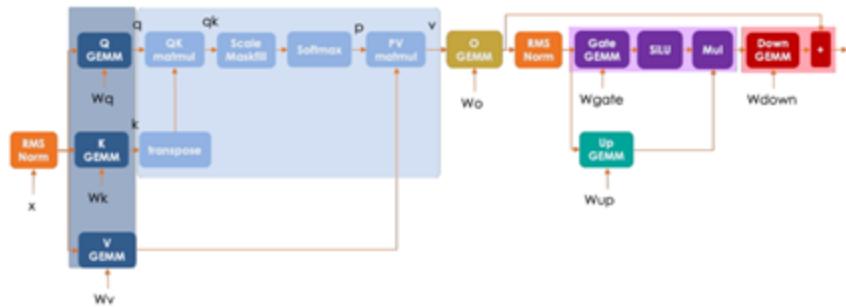
- Native CUDA support in PyTorch
- Leverage CuBLAS, CuDNN, etc.
- Performance tuning can be messy
  - Kernel fusion

Example: Llama3.1 8B



# Section-by-section

- These models can be executed as ***DataFlow*** graphs.
- Break down the model into sections
  - Fewer kernel calls
- Generate output tokens as the data flows through the ***Directed Acyclic Graph (DAG)***
  - Model FLOP utilization (MFU) ↑
  - Model Bandwidth utilization (MBU) ↑
  - Total Cost of ownership ↓



# Need of the hour

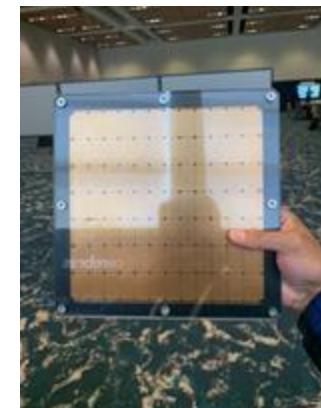
- **Reconfigurable ML accelerators**
  - Field Programmable Gate Array (FPGA)
    - Programmable at fine-grain granularity (Lookup tables, Flip-flops, etc.) - very flexible
    - *Programming is hard!* 
  - Coarse Grain Reconfigurable Architecture (CGRA)
    - Programmable at coarse-grain granularity (ALUs, etc.)
      - flexible enough to lay out transformer DAG on chip
    - *Programming is not so hard!*

# Rise of DataFlow-based ML accelerators

- SambaNova's SN40L Reconfigurable DataFlow Unit (RDU)
  - 1,040 compute cores
  - 3-tier memory within a chip
    - 512 MB of programmable memory (SRAM) - several PB/s
    - 64 GB of HBM - 25+ TB/s
    - 1.5TB of DDR (device memory) - 1600 GB/s
- Groq's Language Processing Units (LPU)
  - Large SRAM ~250 MB per chip
- Cerebras Wafer-scale engine (WSE-3)
  - 900k compute cores
  - 44 GB of SRAM



SambaNova SN40L



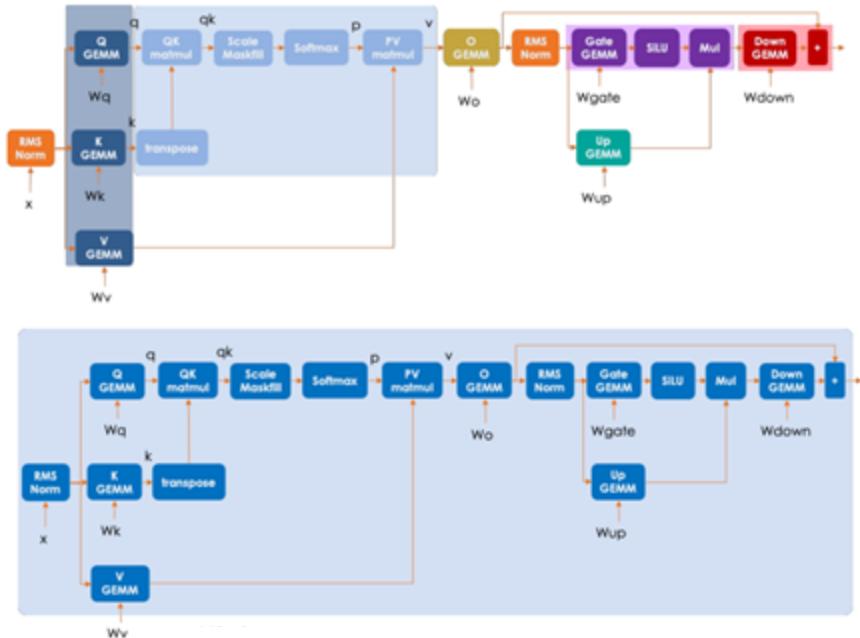
Cerebras WSE-3

# Challenge 3: Extracting the full throughput

# Problems with DataFlow executions

- **High Kernel launch overheads**
  - Reconfigure after each section
- **Low Data Locality**
  - After section is executed
    - Drain data & arguments out of chip
  - Before start of next section
    - Load data & arguments back to chip

Ideal scenario - Execute as a single section



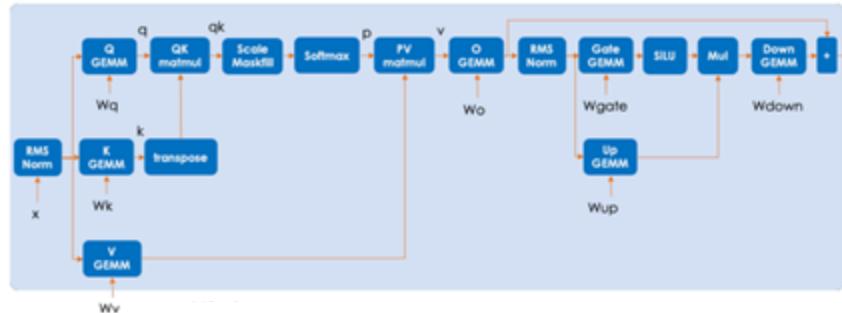
# Inference as DataFlow

- Single kernel launch overhead
- Exploit data locality
- (Recap) Llama 3.1 8B inference

Memory required = 1TB

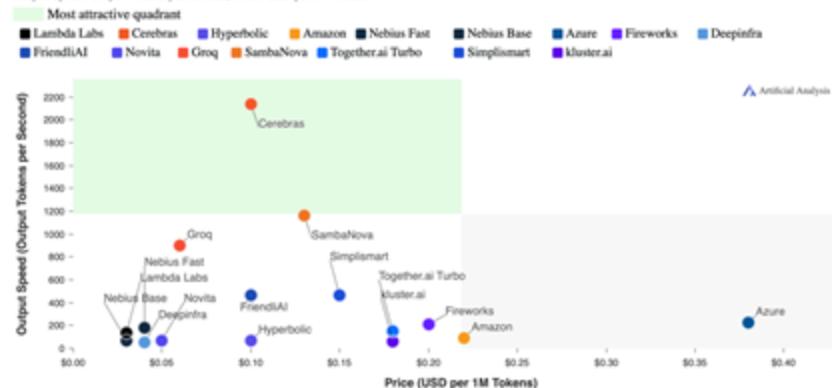
Model FLOPs = ~8.8 TFLOPs

- Can it fit on a single device?
  - Mostly no!
  - Group devices to work in tandem.
- Roofline of single section execution
  - 1000+ tokens/s output speed🔥



Output Speed vs. Price: Llama 3.1 8B Providers

Output Speed: Output Tokens per Second; Price: USD per 1M Tokens

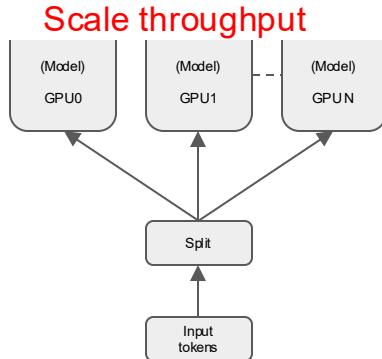


Source: [Artificial analysis](#)

# Challenge 4: Scalability

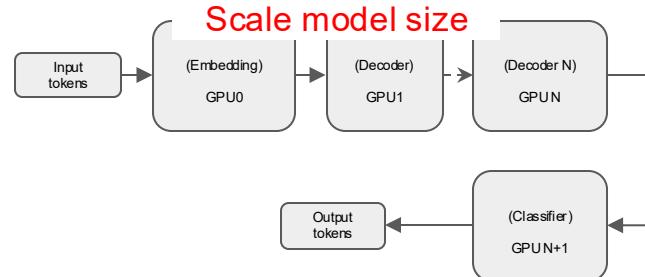
# Data Parallel

- Multiple replicas of the model across devices
  - Process different inputs (e.g., batches of tokens or images)
  - Must fit in a single device



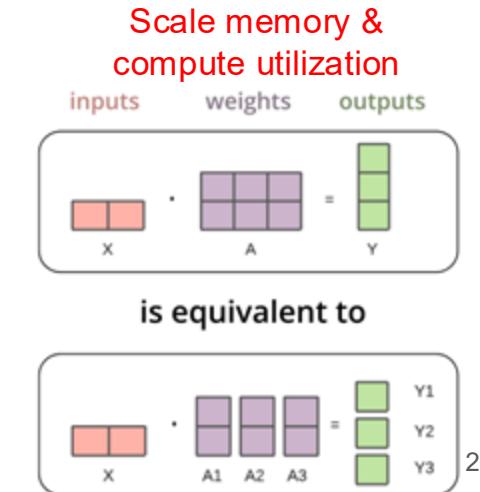
# Model Parallel

- Split the model across devices
  - Pipeline bubbles
  - Not generalizable



# Tensor Parallel

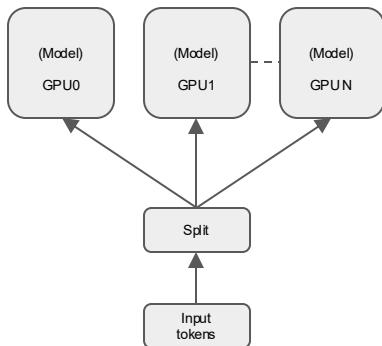
- Multiple replicas of the same operation processing different tensors in parallel
  - Sensitive to communication latency & bandwidth



# Importance of P2P communication

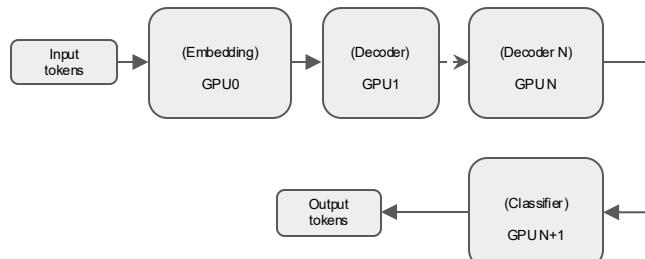
## Data Parallel

- Gradient sync



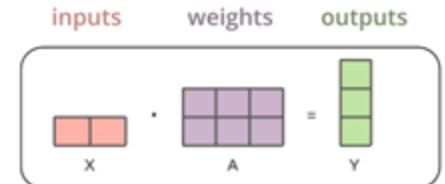
## Model Parallel

- P2P send and receive

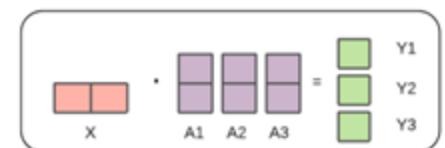


## Tensor Parallel

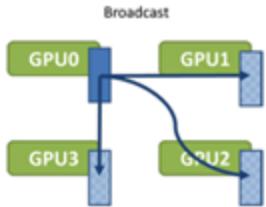
- Many types of sync
- Collective communication calls



is equivalent to

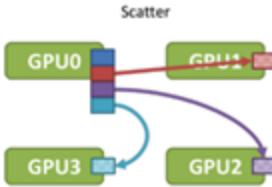


# Collective communication calls & scenarios



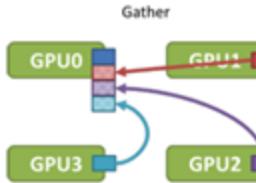
## DP/TP

- Replicate Model params/intermediate results/shared input across devices



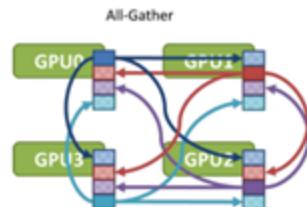
## TP

- Distribute intermediate checkpoints/activations across devices



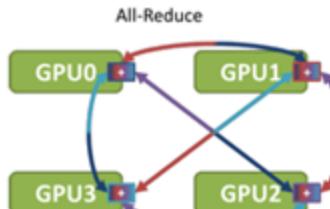
## TP

- Gather outputs after TP op for consecutive single-chip ops (like Norm etc.)



## TP

- Gather outputs after TP op like Embedding.

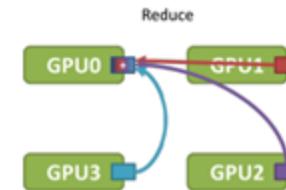


## DP

- Gradient sync

## TP

- Result of 2d-sharded GeMMs



## TP

- Cross Entropy

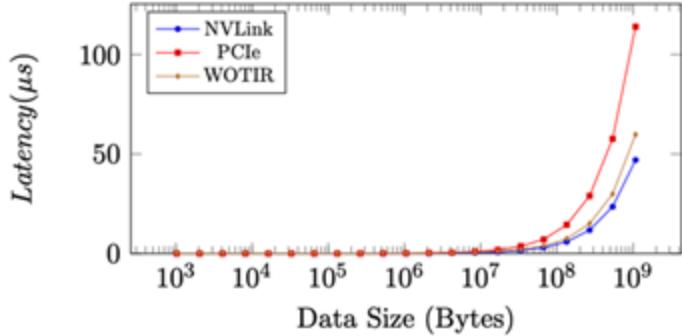
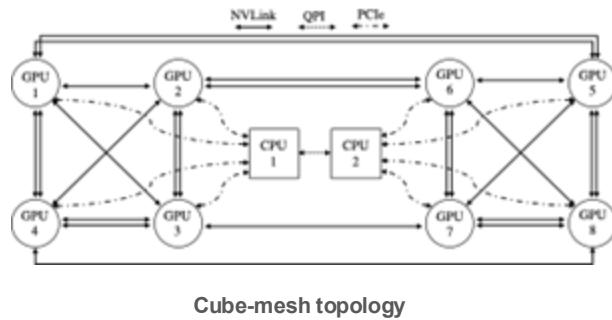
# Need of the hour

- **High-bandwidth communication (p2p)**
  - PCIe
  - NVLink
  - CXL etc.
- **HW/SW co-design**
  - Runtime CCL optimization

# Challenge 5: Designing Low-power systems

# System topology

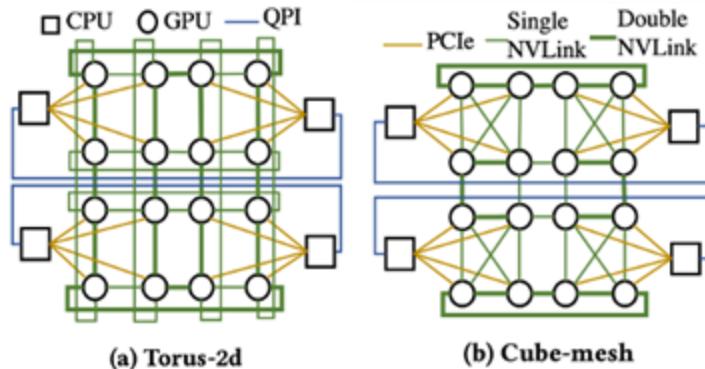
- 8x device cube mesh topology is the common deployment
- Specialized P2P links
  - Nvidia NVlink
  - AMD Infinity fabric
  - CXL
  - PCIe
- Each device can support up to 6 links
- Scaling CCL without A2A connectivity is crucial<sup>[1]</sup>
  - Low-power
  - Low TCO



[1] K. Ranganath, A. Abdolrashidi, S. L. Song and D. Wong, "Speeding up Collective Communications Through Inter-GPU Re-Routing," in IEEE Computer Architecture Letters, vol. 18, no. 2, pp. 128-131, 1 July-Dec. 2019, doi: 10.1109/LCA.2019.2933842. keywords: {Graphics processing units;Bandwidth;Routing;Machine learning;Servers;Training data;Interference;Collective communication;GPU;interconnect}.

# Scaling systems beyond 8 devices

- Specialized interconnects
  - Nvidia NVswitch
    - A2A connectivity up to 16 devices/node
  - High power consumption
  - Low TCO
- Efficient scaling with PCIe
  - Torus
  - Cube-mesh
  - Low-energy
  - Near to linear scaling<sup>[1][2]</sup>
  - Higher TCO



(a) Torus-2d

(b) Cube-mesh

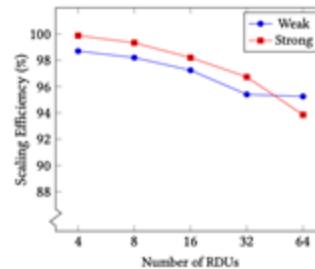


Figure 10: Scaling efficiency on SN30 and RDARuntime

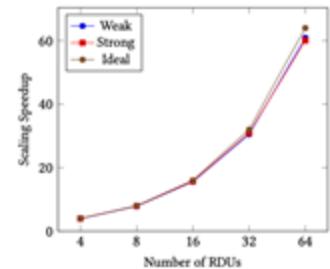


Figure 11: Scaling speedup on SN30 and RDARuntime

[1] Kiran Ranganath, Joshua D. Suetterlein, Joseph B. Manzano, Shuaiwen Leon Song, and Daniel Wong. 2021. MAPA: multi-accelerator pattern allocation policy for multi-tenant GPU servers. In Proceedings of the International Conference for High Performance Computing, Networking, Storage and Analysis (SC '21). Association for Computing Machinery, New York, NY, USA, Article 99, 1–14. <https://doi.org/10.1145/3458817.3480853>

[2] Benjamin Glick, Arjun Sabnis, Renate Kempf, Arnav Goel, Aarti Lalwani, Guoyao Feng, and Kiran Ranganath. 2023. RDARuntime: An OS for AI Accelerators. In Proceedings of the SC '23 Workshops of the International Conference on High Performance Computing, Network, Storage, and Analysis (SC-W '23). Association for Computing Machinery, New York, NY, USA, 1576–1587. <https://doi.org/10.1145/3624062.3624235>

# Challenge 6: DeepSeek, Llama 4 and next?

# Mixture of experts (MoE)

- The success of DeepSeek and Llama 4 demonstrated the benefits of MoE
- Monolithic models are slow and prone to hallucinations
- Benefits of MoE
  - Specialized experts for tasks, languages, etc.
  - Activate only a “subset” of experts for an input token
  - Scaling up to billions of parameters with fixed compute cost
  - Better accuracy than monolithic models
  - Sparsity

# Summary of challenges & HPCA opportunities

<b>Challenges</b>	<b>HPCA search space</b>
1. ML as an HPC problem	Algorithms, roofline analysis
2. Thinking beyond GPUs	Hardware/software architecture
3. Extract the full-throughput	Compiler, roofline analysis, architecture, runtime
4. Scalability	Hardware-software codesign, runtime
5. Design low-power systems	Hardware-software codesign, runtime
6. Mixture of Experts	Algorithms, roofline analysis, Hardware-software codesign, runtime

Thank you!